

AI Alignment as Normative Institutional Design

Maty Bohacek, ??, Roma Patel

Technical approaches to AI alignment have increasingly drawn on concepts developed for problems in moral philosophy, political theory and legal systems—from constitutions and social choice, to deliberation and virtue ethics. However, contemporary technical instantiations of these frameworks bring in these concepts unevenly. Some traditions (e.g., constitutions) have inspired substantial technical AI frameworks, while others remain largely unexplored. Moreover, alignment frameworks often retain only a fraction of the institutional structures developed in the source concept. In order to provide a systematic understanding of this landscape, we construct a two-dimensional taxonomy of normative AI alignment and governance regimes following the typology-construction of [Collier et al. \(2012\)](#). We then provide a mapping of each approach to relevant classical and contemporary philosophy, political theory, and legal theory from which it draws inspiration. We use this framework to understand how normative concepts have so far been translated into technical systems; analyzing what is preserved, what is discarded, and when these omissions represent necessary simplifications rather than design choices. We highlight a recurring tendency of alignment methods to treat normative disagreement as a problem to be *resolved away*, despite governance frameworks that argue for not eliminating it. Lastly, by looking at the entire landscape, we identify missing components and underexplored regimes that could form future directions for building AI systems that can preserve, contest, and govern plurality rather than collapsing it.

Keywords: AI alignment, Governance, Philosophy, Political Theory, Legal Theory, Pluralism.

1. Introduction

[roma: todo citations] What should an AI system do when the values it is expected to follow conflict? How should human preferences that differ across individuals determine its behavior? Should those preferences be aggregated into an objective, encoded as explicit principles, inferred from past decisions, or resolved through deliberation? And when reasonable people continue to disagree, should an aligned system resolve it, or instead preserve and govern that disagreement over time?

These questions are increasingly central to AI alignment, but they are not uniquely questions about AI. Sub-fields of philosophy have long studied how competing values and obligations should be weighed ([Berlin, 1969](#); [Ross, 1930](#)); political theory has asked how institutions can legitimately govern populations that disagree about what is good or just ([Habermas, 1996](#); [Rawls, 1993](#)); and legal jurisprudence has developed systems for specifying, interpreting and contesting

norms over time ([Dworkin, 1988](#); [Hart, 1961](#)). Solutions from these theories have been implemented across institutions: from codified rules and precedent to social convention, voting and pluralism; resulting in fairer institutional regimes for operating and governing under disagreement. In recent years, AI alignment has begun to draw from this repertoire of normative concepts. For example, constitutional AI ([Bai et al., 2022](#)) encodes normative principles in an explicit written artifact that trained AI systems are expected to follow; social-choice approaches to preference learning aggregate judgments according to individuals or incentives ([Ge et al., 2024](#); [Maura-Rivero et al., 2025](#)); deliberative approaches derive principles through collective procedures ([Huang et al., 2024](#)); and pluralistic systems that seek to accommodate heterogeneous and sometimes incompatible values ([Feng et al., 2024](#); [Sorensen et al., 2025](#)). And as a result of the success of these technical approaches, such concepts originating in moral philosophy, jurisprudence and political theory have become a regular part of the technical

vocabulary of alignment.

Yet this transfer of concepts is often highly selective. For example, an alignment system inspired by constitutions such as [Bai et al. \(2022\)](#) may borrow the language of constitutionalism while implementing only a fixed textual specification—it preserves the idea of an authoritative set of principles while omitting the institutional mechanisms through which constitutions are interpreted, contested and revised in judicial systems. Such simplifications are often reasonable engineering choices: technical systems need not exactly reproduce their intellectual sources in every respect. However, our concern is that the **normative properties of a system can change when a concept is separated from the broader institutional structure in which that concept originally played a role**. We call attention to this phenomenon as a problem of *normative translation*. When normative concepts are translated into computational mechanisms with omissions in translation—what consequences follow from these omissions? And if only a subset of a normative tradition is implemented, does this result in a fundamentally different normative concept? We urge such changes in translation should be specified and understood as specific design choices made by practitioners, rather than treated as direct translations of the original theory, as often implied in publications.

[roma: todo rewrite and shorten below! keep introduction at 1.5 pages total?] In order to critically study this, we attempt to systematically map normative alignment and governance approaches in this paper. Rather than analyzing the literature according to the technical alignment mechanisms e.g., preference optimization or constitutional prompting, we ask what normative regime each approach instantiations. Specifically, following the typology from ?, we identify two thematic dimensions that recur across philosophical and legal approaches to normative decision making. The first is about **location and form of normative authority**: specifically, whether the norms that govern a decision exist as a written down specification, a set of examples or cases, a learned disposition, or a decision procedure. The second is about the **disposition of normative con-**

flict: whether norms that conflict during a procedure are collapsed into a single outcome, left open to contestation, or kept as distinct normative orders to be resolved by a downstream process. This 5x3 typology yields a common space for comparing existing alignment and governance methods, as well as a framework to trace the philosophical, political, and legal concepts on which the authors rely, either explicitly or implicitly. This allows us to understand not only which normative concepts have influenced and been instantiated in technical alignment systems, but also to understand the missing gaps and what concepts are yet to be instantiated. Our analysis thus reveals an important property of the current alignment landscape—that the normative design space is unevenly explored. In particular, technical approaches that preserve disagreement rather than collapse it, are less developed. We acknowledge the technical and infrastructural limitations that are a fundamental cause of this, however, we urge researchers to invest in further investigation in these regions could lead to future breakthroughs in alignment.

Our contribution here is therefore threefold. First, we provide a survey and taxonomy of normative concepts that are implemented in both alignment and governance regimes. We provide a searchable encyclopedia of these defined concepts to aid researchers in work involving these concepts. Second, we map existing existing alignment methods to the philosophical concepts they draw from, and construct a 5x3 typology that organizes these concepts along dimensions of normative authority and conflict resolution. Third, we analyze mistranslated concepts in AI alignment systems and consequences of this. We identify underexplored normative regimes and derive design questions for future alignment systems that borrow from the full spectrum of normative concepts. Our goal is not to critique AI systems for not reproducing normative concepts identically, but rather, to enforce that borrowing of normative concepts from traditions should make institutional assumptions visible. By understanding the space of concepts holistically, and identifying the missing conceptual gaps in current alignment systems, we hope to identify both the strengths and limitations of existing approaches, as well as pro-

vide a systematic basis for exploring alternatives for future systems.

2. Related Work

[roma: maybe we move a lot of this to an appendix, and keep half a page worth of the most relevant related work.] A growing body of technical work in AI alignment and governance has been importing concepts from philosophical, legal, and political theory. We first provide an overview of some prominent works, all of which are later included in the constructed mapping of the field. Second, we review existing approaches to taxonomize AI alignment and governance, identify the limitations of their typology, and proceed to describe typology construction mechanisms suitable for their resolution.

Philosophy in AI Alignment. In terms of philosophical concepts, prominent recent examples include work that borrows from value pluralism and incommensurability (Berlin, 2014) as the basis for a roadmap to pluralistic alignment (Sorensen et al., 2024), and work that borrows from accounts of moral patienthood and personal identity (Parfit, 1987; Warren, 1997) to develop methodologies for AI welfare (Long et al., 2024). Other notable examples include the incorporation of paternalism and autonomy (Long et al., 2024) into work on corrigibility (Soares et al., 2015), and the use of virtue ethics and character framing (Anscombe, 2020; Aristotle) in the development of Claude’s Constitution (Anthropic, a,b).

Political Theory in AI Alignment. In terms of political concepts, methods have been developed on the basis of social contract theory (Hobbes, 1885; Locke; Rahwan, 2018; Rousseau). Others borrow from social choice theory and preference aggregation (Arrow, 2012; Sen, 2017) in developing methods that avoid a "tyranny of the majority" (Bao et al., 2026). Others employ the veil of ignorance (Rawls, 2017) to align systems with principles of justice (Weidinger et al., 2023). Concepts from deliberative democracy (Cohen, 2005; Habermas, 2015) have been imported into

a number of works, including efforts to source principles from public input (Huang et al., 2024) and to grade how democratically a given process is run (Ovadya et al., 2024). In parallel, constitutionalist concepts have been employed in the development of AI constitutions (Bai et al., 2022).

Legal Theory in AI Alignment. Finally, in terms of legal concepts, notable examples of inspiration from classical theoretical work include the use of principle-oriented interpretivism (Dworkin, 1988) and of a positivist account of law as analogical reasoning (Hart, 2012; Sunstein, 2018), both mapped onto Constitutional AI and case-based reasoning (Caputo, 2024). This work also employs concepts such as precedent, stare decisis (i.e., doctrine instructing the system to stand by things already decided), and case-based reasoning from common law doctrine (Llewellyn, 2016). These concepts are also formative for proposals that pair a constitution with AI courts developing AI case law (Abiri, 2024). Statutory and constitutional interpretation methods (Scalia and Garner, 2012) are also important and have been adapted as tools for resolving ambiguity in specifications (Kolt et al., 2026).

Taxonomies of Alignment Approaches. Existing surveys of alignment and governance work generally organize the field by technical methods such as RLHF (Lu et al., 2025; Shen et al., 2023), or by value-setting principles (Chalkidis, 2025). For example, one line of work distinguishes between Overton, steerable, and distributional pluralism (Sorensen et al., 2024). Others have organized the field into pluralistic, contextual, and collective mechanisms (Bao et al., 2026). Still, the field remains largely unmapped, with the various approaches emerging largely independently of each other and not in direct conversation with one another.

Typology Construction. Motivated by taxonomies sorted by non-equivalent criteria and by free-floating typologies whose dimensions are never made explicit, prior work provides a template for building typologies from an overarching concept into explicit, clearly defined dimensions

of an orthogonal nature (Collier et al., 2012). The underlying notion of a property space is due to earlier work in the same tradition (Barton, 1955), as is the distinction between conceptual and explanatory typologies (Elman, 2005). As we find equivalent failure types in existing taxonomies of alignment, we follow this methodology.

Comparison to Related Work. Existing work maps traditional concepts from philosophical, political, and legal theory onto alignment and governance methods one at a time. We take a different approach by organizing the field according to a more complete set of underlying principles by which these methods organize themselves, and map individual traditions and concepts onto them subsequently. Furthermore, we attempt to critically interrogate whether the stated inspiration in traditional work results in a faithful technical implementation, and describe ways in which gaps or confusion emerge. Additionally, this approach allows us to identify gaps in approaches to alignment and governance which remain largely unexplored in the technical literature.

3. A Normative Design Space for Alignment: A Taxonomy

[roma: this needs to be the biggest section. we need to clarify how the taxonomy was created/why this current form is the best]

In order to study the use of normative concepts across fields, we create a two-dimensional property space defined by two axes: the locus of normative authority and the disposition of normative conflict.

We derive these dimensions through a systematic mapping of alignment and governance literature followed by conceptual substruction: overlapping surface categories are decomposed into the underlying attributes that distinguish them. Following Petersen et al. (2008), two coders independently extract normative concepts from the sampled literature and reconcile near-equivalent categories; full collection and inclusion criteria appear in Appendix A. Unlike conventional typologies that discard unrealized combinations,

we retain empty cells because absence from the technical literature is itself analytically informative for identifying underexplored architectures.

We therefore follow prior work in using substruction (Barton, 1955) to recover underlying dimensions and extend them into a full property space here (Lazarsfeld, 1937). Substruction is "the procedure of finding, for a given system of types, the attribute space in which it belongs and the reduction which has been implicitly used" (Lazarsfeld, 1937). It is appropriate here because the categories which we previously derived overlap as projections of a higher-dimensional space onto a single axis. We deem this procedure applicable since prior work also considers normative types and arrives at axes similar to ours (Barton, 1955), as shown later.

This process yields two axes: locus of normative authority, which indicates where the binding normative content resides prior to the moment of decision, which here corresponds to model inference, and disposition of normative conflict, which indicates how any conflicts should be resolved. Their cross product gives a 5 x 3 property space. We note a difference from the null resolution procedure of prior work (Lazarsfeld, 1937). Whereas that procedure removes any null cells, we retain them, since we are interested in gaps in the literature.

For the locus of authority axis, we find the following categories:

1. **Codified.** The norms the AI ought to follow exist as an explicit written specification, set prior to deployment and, in principle, complete before any particular case arises. Authority resides in this specification (Hart, 2012).
2. **Case-accumulated.** The norms reside in a growing body of past decided cases. No case is decided by a rule stated in advance; what the AI should do is whatever the accumulated decisions have held (Jonsen and Toulmin, 1988; Levi, 2013).
3. **Customary.** The norms reside in regularities of practice (what agents in the population actually do), never written down and never adjudicated. There is no authoritative moment

at which a norm gets explicitly stated (Burke, 1986; Lewis, 2008).

4. **Dispositional.** The norms reside in the AI's trained character rather than in any external artifact. There is no rule to consult and nothing to look up, only the disposition itself (Aristotle).
5. **Procedural.** No norms persist between decisions. What the AI should do is whatever a specified procedure (e.g., a vote, a deliberation, a debate) yields at the moment of decision (Habermas, 2015).

For the disposition of normative conflict, we find the following categories:

1. **Collapse.** When the AI's norms conflict, the system resolves the conflict into a single output, without trace or acknowledgment of the disagreement that produced it. The plurality is thereby aggregated, averaged, or adjudicated away (Sorensen et al., 2024).
2. **Contest.** The conflict between norms is preserved and kept procedurally live. Either the output itself carries the disagreement, or the decision stays open to challenge through a standing channel (Mouffe, 2000; Pettit, 1997).
3. **Partition.** The conflict is neither resolved nor sustained but routed around. Multiple normative orders are kept intact, and a rule assigns which one governs a given case (de Sousa Santos, 2002; Lijphart, 1977; Savigny, 1867).

We considered open, axial and selective coding (Corbin and Strauss, 2014), as well as the first-order, second-order and aggregate-dimension ladder (Gioia et al., 2013), which would perhaps read as more familiar to a social science reader. We prefer the keyword-based approach with dimension resolution, as it is more native to the computer science literature it addresses. However, we note that the alternatives would likely yield a comparable outcome in terms of the intermediate categories.

General morphological analysis (Ritchey, 2022) would be inappropriate here, given that it con-

structs its multi-dimensional space while inspecting unrealized combinations, which requires preemptive stipulation. Since we are working with an existing structure and literature landscape, this would introduce unnecessary assumptions and likely lead to a less grounded result.

4. Case Studies in the Design Space

[roma: this can cover 2-3 case studies in about 1 page.] We use our 53 typology to examine a few different alignment methods that are widely used in the literature. In each case study we assess the fidelity of the implementation of the method as compared to its stated theoretical roots. We categorize this fidelity (e.g., whether it was a computational or engineering constraint)

5. AI Instantiations of Political Theory

5.1. Reinforcement Learning from Human Feedback (RLHF)

RLHF as a post-training paradigm is now widely used in all LLM-based AI products as a means of aligning them to user preferences. This method trains a reward model by asking human raters to select preferred responses over lots of samples of different text responses. This data is then used to fit the reward model, through a scoring function that can then predict human preferences over unseen text responses. This model is then used within a reinforcement learning paradigm with optimization algorithms (e.g., PPO) to provide the approximated human preference rewards to an LLM in order to steer it towards maximizing human preferences.

This idea is rooted in Utilitarianism: Bentham and Mill posit that an ideal decision is one that maximizes net utility across a population of users. RLHF as an algorithm does exactly this. Assuming that a population of user preferences and values can be collapsed into a single continuous scalar metric, it treats human preferences as homogeneous, and something can be averaged at a population level, while ignoring individual user preferences that might contradict others. Critiques against this include Tocqueville and Mill

that warn against the "tyranny of the majority" i.e., a resulting state where majoritarian consensus can marginalize, erase or suppress minority perspectives and contradicting world views. In the context of RLHF, marginalized populations of users that are less represented—or even not included at all—in the rater distributions will have their preferences and values suppressed under current training algorithms. This disallows for distributional pluralism and only optimizes for the majoritarian consensus. Separately, Social Choice Theory provides theoretical arguments against the validity of these frameworks. Arrow's Impossibility Theorem guarantees that no rank-order voting system can convert individual preferences into a population-level consensus while maintaining fairness axioms. Hence, whenever humans have conflicting preferences, RLHF-like algorithms with preference aggregation result in impossible states. Individual intransitivity compounds this further. These algorithms attempt to solve through gradient descent what political philosophy has proved cannot truly be solved without political trade-offs and suppression of certain viewpoints. Lastly, because the collection of human preferences and values are dictated by the curators e.g., an AI lab that uses specific vendor guidelines, rater demographics to include, red-teaming rubrics for data; this enforces a system similar to a Hobbesian sovereign, where the boundary conditions of acceptable feedback decide what demographics or populations of people to consult, thus enforcing a top-down normative centralization under the guise of reflective general and unbiased human feedback.

These algorithms therefore forego Mouffean agonistic pluralism—instead of preserving conflict and differing worldviews, it collapses disagreement and preferences into a safer aggregate of worldviews. It also collects passive preference data, rather than engaging in a Habermasian deliberative framework that allows for users to partake in an active, deliberative interaction to reason with.

5.2. Constitutional AI (CAI)

Moving away from utilitarian preference aggregation, CAI turns to statutory constitutionalism and

legal positivism, which sees law as a system of explicit rules created by a recognized authority, whose legitimacy stems from the formal correctness and enactment of said rules, not necessarily a consensus from involved parties. The CAI framework embodies this with a written "Constitution". This takes form in a system prompt or model spec, guiding the training of a model through self-critique of its actions against the constitution, iterative revision of constitutions, and fine-tuning of the model based on this whole process. It can also be used to score the quality of the model's behavior. In addition to positivism, this also embodies Fuller's Morality of Law, since a model's alignment is, at least in theory, human-understandable and statutory: unlike with the hardly interpretable preference vectors in RLHF, the constitution, as a document in natural language, contains easily understandable and auditable principles, which can be inspected and updated over time.

However, this approach also foregoes multiple constructs important to the original theory upon which it draws. Under this approach, the constitution document is usually created and edited internally at the respective AI lab which owns the model, and is later imposed universally on users without any democratic ratification while lacking a recourse mechanisms. This top-down, paternalistic enforcement diverges from X and Y. If we look to democratic systems, where written constitutions typically derive their authority from participatory deliberation or elected representations, absence of such guarantees would certainly undermine the overall legitimacy of such a governance system.

Moreover, since users are treated as subjects, whose agreement with the written constitution is taken for granted (often, the constitution itself is not even available for inspection, with the notable exception of Claude's constitution), this approach also assumes Lockean consent. This is contrary to treating users as active citizens who could participate in drafting and amending of these constitutions if they so chose, which the X/Y PRINCIPLE would otherwise invite. Furthermore, CAI embodies an implicit assumption that complex, context-dependent moral dilemmas can

be anticipated and resolves through static principles, foregoing a case-by-case and bottom-up reasoning. Finally, the principle of Pettitian contestability or recourse is also violated here, since the CAI protocol, as it is commonly implemented, does not create any appeal mechanisms, where users could contest the application of a constitutional rule in practice. This violation holds even in cases where the constitution document is available for inspection.

Notably, over time, the term "Constitution" has become overloaded and the discourse surrounding CAI in general has become confused. In its original meaning, the constitution was statutory. Similar to legal positivism, it derived its authority from enactment by a governing body, which was recognized and often implicitly assumed to be democratic. Here, the constitution document was used to adjudicate conflicts between competing values and ultimately measure model quality, as suggested above. In its second, more recent meaning, the constitution has assumed a more constitutive role, closer to Aristotelian virtue ethics, with the assumption that correct actions follow from a well-cultivated character. Here, the constitution instantiates a character by defining its general dispositions instead of stating rules.

Critically, these two meanings engage in distinct forms of normative work. The statutory meaning enjoys the procedural legitimacy of constitutional government, which is something a privately created rulebook, on its face, would not. Meanwhile, the constitutive meaning enjoys behavioral coherence as well as relational trust, with the assumption that a model with a trustworthy character can be reasonably expected to generalize well to unseen cases, hard to preempt in a traditional, statutory-oriented constitution. The latter meaning also invites the trust ordinarily extended to persons, which policy frameworks would not.

The dominant way in which CAI is temporarily deployed is best characterized as a combination of the two meanings, as they are not mutually exclusive in terms of the written document (which can combine general character traits as well as rules and restrictions) nor technical implementation (the document can still be used during

training and inference monitoring). The slippage between the two, however, does not resolve the deficits outlined above. Pettitian contestability becomes harder to maintain under such a combination, because dispositions, unlike specific rules, offer no textual criterion to be used in cases where the user wishes to contest specific model behavior. Virtue ethics, on the other hand, presupposes a community whose practices cultivate and ratify the sought virtues, which these constitutions inherently violate by being authored and maintained internally within the labs.

5.3. Adversarial Red-teaming

"Red-teaming" frameworks have been applied in LLM safety pipelines to employ adversarial tactics in a systematic and automated manner, in order to find hidden vulnerabilities, unsafe behaviors and edge-case failures before models are deployed. These methods range from curated prompt sets, human-in-the-loop methods, and LLM-based multi-agent training pipelines that generate adversarial data for the AI system in question. This framework maps directly to Madisonian institutional design positing that robust governance systems do not rely solely on written codifications and good intentions, but require active, institutionalized opposition to test the resilience of state against edge-case failure models, corruption or subversion. Separately, Popper argued that open societies remain resilient not just by verifying and updating their norms, but by subjecting them to criticisms and empirical falsification. Moreover, agnostic frameworks view adversarial red-teaming as an essential and ongoing negotiative procedure in a system that map the true operational boundaries of the model's safety space.

While red-teaming is a crucial component of AI systems, the way that it is implemented in as-is lacks important aspects that these political theories outlined. For one, red-teaming is usually treated as a bounded, pre-deployment question-answering task rather than a publicly available civic mechanism. It therefore foregoes Habermasian deliberative intent, by only being conducted under corporate contractors, thus foregoing democratic transparency and public par-

ticipation. It also does not allow for Pettition contestatory recourse, when only used as a static, pre-deployment filter.

While outside of the current alignment and red-teaming pipelines, users could be empowered to challenge model behavior after deployment through dynamic appeal channels or live auditing interfaces. Lastly, red-teaming rubrics and adversarial training regimes have so far been defined almost exclusively for legal liability, PR safety, and other internally-defined risks, and hence resided within corporate risk-compliance frameworks. Any minority-group harms that fall outside these corporate incentives of safety or broader, publicly held viewpoints of safety are disregarded by these approaches.

5.4. Automated Critics and RLAIIF

5.5. Pluralistic and Value Alignment

5.6. Collective CAI

6. Expanding the Design Space

[roma: this can discuss what is missing and future gaps to fill]

7. Conclusions

[roma: also either bleed in limitations here, or have a section for it before this]

References

G. Abiri. Public constitutional AI. *Peking University School of Transnational Law Research Paper, Georgia Law Review*, 59, 2024.

G. E. Anscombe. Modern moral philosophy. In *The definition of morality*, pages 211–234. Routledge, 2020.

Anthropic. Claude’s constitution, a. URL <https://www.anthropic.com/constitution>.

Anthropic. Claude’s new constitution, b. URL <https://www.anthropic.com/news/claude-new-constitution>.

Aristotle. *Nicomachean Ethics*, volume 9 of *The Works of Aristotle Translated into English*. Clarendon Press.

K. J. Arrow. *Social choice and individual values*, volume 12. Yale university press, 2012.

Y. Bai et al. Constitutional AI: Harmlessness from AI feedback, 2022.

H. Bao, Y. Huang, X. Wang, Z. Zhang, Y. Zhou, C. Yang, X. Zhang, and Y. Ye. AI alignment breaks at the edge. *arXiv preprint arXiv:2602.20042*, 2026.

A. H. Barton. The concept of property-space in social research. *The language of social research*, 111:40, 1955.

I. Berlin. *Four Essays on Liberty*. Oxford University Press, London and New York, 1969.

I. Berlin. Two concepts of liberty. In *Reading political philosophy*, pages 231–237. Routledge, 2014.

E. Burke. Reflections on the revolution in france (1790). *Everyman’s Library*, 1986.

N. A. Caputo. Alignment as jurisprudence. *Yale Journal of Law and Technology (forthcoming)*, 2024.

I. Chalkidis. Decoding alignment: A critical survey of LLM development initiatives through value-setting and data-centric lens. *arXiv preprint arXiv:2508.16982*, 2025.

J. Cohen. Deliberation and democratic legitimacy. In *Debates in contemporary political philosophy*, pages 352–370. Routledge, 2005.

D. Collier, J. LaPorte, and J. Seawright. Putting typologies to work: Concept formation, measurement, and analytic rigor. *Political research quarterly*, 65(1):217–232, 2012.

J. Corbin and A. Strauss. *Basics of qualitative research: Techniques and procedures for developing grounded theory*. Sage publications, 2014.

B. de Sousa Santos. Toward a new legal common sense: Law. *Globalization and Emancipation*, 314, 2002.

- R. Dworkin. *Law's empire*. Harvard university press, 1988.
- C. Elman. Explanatory typologies in qualitative studies of international politics. *International Organization*, 59(2):293–326, 2005. doi: 10.1017/S0020818305050101.
- S. Feng, T. Sorensen, Y. Liu, J. Fisher, C. Y. Park, Y. Choi, and Y. Tsvetkov. Modular pluralism: Pluralistic alignment via multi-LLM collaboration. In Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 4151–4171, Miami, Florida, USA, Nov. 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.240. URL <https://aclanthology.org/2024.emnlp-main.240/>.
- L. Ge, D. Halpern, E. Micha, A. D. Procaccia, I. Shapira, Y. Vorobeychik, and J. Wu. Axioms for ai alignment from human feedback. *Advances in Neural Information Processing Systems*, 37:80439–80465, 2024.
- D. A. Gioia, K. G. Corley, and A. L. Hamilton. Seeking qualitative rigor in inductive research: Notes on the gioia methodology. *Organizational research methods*, 16(1):15–31, 2013.
- J. Habermas. *Between Facts and Norms: Contributions to a Discourse Theory of Law and Democracy*. MIT Press, Cambridge, MA, 1996.
- J. Habermas. *Between facts and norms: Contributions to a discourse theory of law and democracy*. John Wiley & Sons, 2015.
- H. L. A. Hart. *The Concept of Law*. Clarendon Press, Oxford, 1961.
- H. L. A. Hart. *The concept of law*. OUP Oxford, 2012.
- T. Hobbes. *Leviathan, or, The matter, forme, & power of a common-wealth ecclesiasticall and civill*, volume 21. 1885.
- S. Huang, D. Siddarth, L. Lovitt, T. I. Liao, E. Durmus, A. Tamkin, and D. Ganguli. Collective constitutional AI: Aligning a language model with public input. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1395–1417, 2024.
- A. R. Jonsen and S. Toulmin. *The abuse of casuistry: A history of moral reasoning*. Univ of California Press, 1988.
- N. Kolt, N. Caputo, J. Boeglin, C. O’Keefe, R. Bommasani, S. Casper, M.-F. Cuéllar, N. Feldman, I. Gabriel, G. K. Hadfield, et al. Legal alignment for safe and ethical AI. *arXiv preprint arXiv:2601.04175*, 2026.
- P. F. Lazarsfeld. Some remarks on the typological procedures in social research. *Zeitschrift für Sozialforschung*, 6(1):119–139, 1937.
- E. H. Levi. *An introduction to legal reasoning*. University of Chicago Press, 2013.
- D. Lewis. *Convention: A philosophical study*. John Wiley & Sons, 2008.
- A. Lijphart. *Democracy in plural societies: A comparative exploration*. Yale University Press, 1977.
- K. N. Llewellyn. *The common law tradition: deciding appeals*, volume 16. Quid Pro Books, 2016.
- J. Locke. *Two Treatises of Government*. Cambridge Texts in the History of Political Thought. Cambridge University Press, student edition.
- R. Long, J. Sebo, P. Butlin, K. Finlinson, K. Fish, J. Harding, J. Pfau, T. Sims, J. Birch, and D. Chalmers. Taking AI welfare seriously. *arXiv preprint arXiv:2411.00986*, 2024.
- H. Lu, L. Fang, R. Zhang, X. Li, J. Cai, H. Cheng, L. Tang, Z. Liu, Z. Sun, T. Wang, et al. Alignment and safety in large language models: Safety mechanisms, training paradigms, and emerging challenges. *arXiv preprint arXiv:2507.19672*, 2025.
- R.-R. Maura-Rivero, C. Nagpal, R. Patel, and F. Visin. Utility-inspired reward transformations improve reinforcement learning training of language models, 2025. URL <https://arxiv.org/abs/2501.06248>.

- C. Mouffe. *The democratic paradox*. verso, 2000.
- A. Ovadya, K. Redman, L. Thorburn, Q. Z. Chen, O. Smith, F. Devine, A. Konya, S. Milli, M. Revel, K. Feng, et al. Democratic AI is possible. the democracy levels framework shows how it might work. *arXiv preprint arXiv:2411.09222*, 2024.
- D. Parfit. *Reasons and persons*. Oxford University Press, 1987.
- K. Petersen, R. Feldt, S. Mujtaba, and M. Mattsson. Systematic mapping studies in software engineering. In *12th international conference on evaluation and assessment in software engineering (EASE)*. BCS Learning & Development, 2008.
- P. Pettit. *Republicanism: A theory of freedom and government*. Oxford University Press, 1997.
- I. Rahwan. Society-in-the-loop: programming the algorithmic social contract. *Ethics and information technology*, 20(1):5–14, 2018.
- J. Rawls. *Political Liberalism*. Columbia University Press, New York, 1993.
- J. Rawls. A theory of justice. In *Applied ethics*, pages 21–29. Routledge, 2017.
- T. Ritchey. General morphological analysis: an overview. 2022.
- W. D. Ross. *The Right and the Good*. Clarendon Press, Oxford, 1930.
- J.-J. Rousseau. *The Social Contract and Other Later Political Writings*. Cambridge Texts in the History of Political Thought. Cambridge University Press.
- F. K. v. Savigny. System of the modern roman law. (*No Title*), 1867.
- A. Scalia and B. A. Garner. *Reading Law: The Interpretation of Legal Texts*. Thomson/West, St. Paul, MN, 2012. ISBN 978-0-314-27555-4.
- A. Sen. *Collective choice and social welfare: Expanded edition*. Penguin UK, 2017.
- T. Shen, R. Jin, Y. Huang, C. Liu, W. Dong, Z. Guo, X. Wu, Y. Liu, and D. Xiong. Large language model alignment: A survey. *arXiv preprint arXiv:2309.15025*, 2023.
- N. Soares, B. Fallenstein, E. Yudkowsky, and S. Armstrong. Corrigibility. In *AAAI Workshops: Workshops at the Twenty-Ninth AAAI Conference on Artificial Intelligence*, Austin, TX, 2015. AAAI Publications.
- T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Mireshghallah, C. M. Rytting, A. Ye, L. Jiang, X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024.
- T. Sorensen, P. Mishra, R. Patel, M. H. Tessler, M. A. Bakker, G. Evans, I. Gabriel, N. Goodman, and V. Rieser. Value profiles for encoding human variation. In C. Christodoulopoulos, T. Chakraborty, C. Rose, and V. Peng, editors, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2047–2095, Suzhou, China, Nov. 2025. Association for Computational Linguistics. ISBN 979-8-89176-332-6. doi: 10.18653/v1/2025.emnlp-main.106. URL <https://aclanthology.org/2025.emnlp-main.106/>.
- C. R. Sunstein. *Legal reasoning and political conflict*. Oxford University Press, 2018.
- M. A. Warren. *Moral status: Obligations to persons and other living things*. Clarendon Press, 1997.
- L. Weidinger, K. R. McKee, R. Everett, S. Huang, T. O. Zhu, M. J. Chadwick, C. Summerfield, and I. Gabriel. Using the veil of ignorance to align AI systems with principles of justice. *Proceedings of the National Academy of Sciences*, 120(18): e2213709120, 2023.

Acknowledgements