

AI Alignment as Normative Institutional Design

Maty Bohacek, ??, Roma Patel

Technical approaches to AI alignment have increasingly drawn on concepts developed for problems in moral philosophy, political theory and legal systems—from constitutions and social choice, to deliberation and virtue ethics. However, contemporary technical instantiations of these frameworks bring in these concepts unevenly. Some traditions (e.g., constitutions) have inspired substantial technical AI frameworks, while others remain largely unexplored. Moreover, alignment frameworks often retain only a fraction of the institutional structures developed in the source concept. In order to provide a systematic understanding of this landscape, we construct a two-dimensional taxonomy of normative AI alignment and governance regimes following the typology-construction of Collier et al. (2012). We then provide a mapping of each approach to relevant classical and contemporary philosophy, political theory, and legal theory from which it draws inspiration. We use this framework to understand how normative concepts have so far been translated into technical systems; analyzing what is preserved, what is discarded, and when these omissions represent necessary simplifications rather than design choices. We highlight a recurring tendency of alignment methods to treat normative disagreement as a problem to be *resolved away*, despite governance frameworks that argue for not eliminating it. Lastly, by looking at the entire landscape, we identify missing components and underexplored regimes that could form future directions for building AI systems that can preserve, contest, and govern plurality rather than collapsing it.

Keywords: AI alignment, Governance, Philosophy, Political Theory, Legal Theory, Pluralism.

1. Introduction

[roma: todo citations] What should an AI system do when the values it is expected to follow conflict? How should human preferences that differ across individuals determine its behavior? Should those preferences be aggregated into an objective, encoded as explicit principles, or resolved through deliberation? And when reasonable people continue to disagree, should an aligned system resolve it at all, or instead preserve and govern that disagreement over time?

These questions are increasingly central to AI alignment, but they are not uniquely questions about AI. Sub-fields of philosophy have long studied how competing values and obligations should be weighed (Berlin, 1969; Ross, 1930); political theory has asked how institutions can legitimately govern populations that disagree about what is good or just (Habermas, 1996; Rawls, 1993); and legal jurisprudence has developed systems for specifying, interpreting and contesting norms over time (Dworkin, 1988; Hart, 1961).

From these theories, real instantiations of concepts have been implemented across institutions: from codified rules in constitutions to democratic procedures aggregating competing preferences; resulting in different institutional regimes for operating and governing under disagreement. In recent years, AI alignment has begun to draw from this repertoire of normative concepts. For example, constitutional AI (Bai et al., 2022) encodes normative principles in an explicit written artifact that trained AI systems are expected to follow; social-choice approaches to preference learning aggregate judgments according to individuals or incentives (Ge et al., 2024; Maura-Rivero et al., 2025); deliberative approaches derive principles through collective procedures (Huang et al., 2024); and pluralistic systems attempt to allow model outputs to represent heterogeneous and sometimes incompatible values (Feng et al., 2024; Sorensen et al., 2025). And as a result of the success of these approaches, such concepts originating in moral philosophy, jurisprudence and political theory have become a regular part of the

technical vocabulary of alignment.

Yet this transfer of concepts is often selective. For example, an alignment system inspired by constitutions (Bai et al., 2022) may borrow the codification of constitutions while implementing only a part of the entire framework—it preserves the idea of an authoritative set of principles written for a system, while omitting the institutional mechanisms through which these principles are interpreted, contested and revised, as they are in judicial systems. Such simplifications are often reasonable engineering choices: technical systems need not exactly reproduce their intellectual sources in every respect due to infrastructural or engineering constraints. However, our broader concern is that the **normative properties of a system can change when a concept is separated from the broader institutional structure in which that concept originally played a role**. We call attention to this phenomenon as a problem of *normative translation*. In this work, rather than asking whether technical implementations faithfully reproduce the theoretical concepts they originated from, we ask which features are preserved in translation, which are omitted, and what consequences follow from these choices.

[roma: todo rewrite and shorten below! keep introduction at 1.5 pages total?] In order to critically study this, we systematically map normative concepts to alignment implementations, by asking what normative regime each alignment method instantiates. Specifically, following Collier et al. (2012)’s typological approach, we identify two recurring thematic dimensions and use these to organize concepts into a multidimensional property space. The first dimension is about **location and form of normative authority**: specifically, whether the norms that govern a decision exist as any of the five sub-dimensions: a written down specification, a set of examples or cases, a set of unwritten practices or customs, a learned outcome or settlement, or a decision procedure. The second is about the **disposition of normative conflict**, categorized into three sub-dimensions: whether norms that conflict during a procedure are collapsed into a single outcome, left open to contestation, or kept as distinct normative orders to be resolved by a downstream

process. This 5x3 typology yields a common space for comparing existing alignment and governance methods, that trace back to philosophical, political, and legal concepts on which the works rely, either explicitly or implicitly. This allows us to understand not only which normative concepts have influenced and been instantiated in technical alignment systems, but also to understand the missing gaps and what concepts are yet to be instantiated. Our analysis thus reveals an important property of the current alignment landscape—that the normative design space is unevenly explored. In particular, technical approaches that preserve disagreement rather than collapse it, are less developed. We acknowledge the technical and infrastructural limitations that are a fundamental cause of this, however, we urge researchers to invest in further investigation in these regions that could lead to future breakthroughs in alignment.

Our contribution here is threefold. First, we provide a survey and taxonomy of normative concepts that are implemented in both alignment and governance regimes. We provide a searchable encyclopedia of these defined concepts to aid researchers in work involving these concepts. Second, we map existing existing alignment methods to the philosophical concepts they draw from, and construct a 5x3 typology that organizes these concepts along dimensions of normative authority and conflict resolution. Third, we analyze mistranslated concepts in AI alignment systems and consequences of this. We identify underexplored normative regimes and derive design questions for future alignment systems that borrow from the full spectrum of normative concepts. Our goal is not to critique AI systems for not reproducing normative concepts identically, but rather, to enforce that borrowing of normative concepts from traditions should make institutional assumptions visible. By understanding the space of concepts holistically, and identifying the missing conceptual gaps in current alignment systems, we hope to identify both the strengths and limitations of existing approaches, as well as provide a systematic basis for exploring alternatives for future systems.

2. Related Work

[roma: maybe we move a lot of this to an appendix, and keep half a page worth of the most relevant related work.] A growing body of technical work in AI alignment and governance has been importing concepts from philosophical, legal, and political theory. We first provide an overview of some prominent works, all of which are later included in the constructed mapping of the field. Second, we review existing approaches to taxonomize AI alignment and governance, identify the limitations of their typology, and proceed to describe typology construction mechanisms suitable for their resolution.

Philosophy in AI Alignment. In terms of philosophical concepts, prominent recent examples include work that borrows from value pluralism and incommensurability (Berlin, 2014) as the basis for a roadmap to pluralistic alignment (Sorensen et al., 2024), and work that borrows from accounts of moral patienthood and personal identity (Parfit, 1987; Warren, 1997) to develop methodologies for AI welfare (Long et al., 2024). Other notable examples include the incorporation of paternalism and autonomy (Long et al., 2024) into work on corrigibility (Soares et al., 2015), and the use of virtue ethics and character framing (Anscombe, 2020; Aristotle) in the development of Claude’s Constitution (Anthropic, a,b).

Political Theory in AI Alignment. In terms of political concepts, methods have been developed on the basis of social contract theory (Hobbes, 1885; Locke; Rahwan, 2018; Rousseau). Others borrow from social choice theory and preference aggregation (Arrow, 2012; Sen, 2017) in developing methods that avoid a "tyranny of the majority" (Bao et al., 2026). Others employ the veil of ignorance (Rawls, 2017) to align systems with principles of justice (Weidinger et al., 2023). Concepts from deliberative democracy (Cohen, 2005; Habermas, 2015) have been imported into a number of works, including efforts to source principles from public input (Huang et al., 2024) and to grade how democratically a given process

is run (Ovadya et al., 2024). In parallel, constitutionalist concepts have been employed in the development of AI constitutions (Bai et al., 2022).

Legal Theory in AI Alignment. Finally, in terms of legal concepts, notable examples of inspiration from classical theoretical work include the use of principle-oriented interpretivism (Dworkin, 1988) and of a positivist account of law as analogical reasoning (Hart, 2012; Sunstein, 2018), both mapped onto Constitutional AI and case-based reasoning (Caputo, 2024). This work also employs concepts such as precedent, stare decisis (i.e., doctrine instructing the system to stand by things already decided), and case-based reasoning from common law doctrine (Llewellyn, 2016). These concepts are also formative for proposals that pair a constitution with AI courts developing AI case law (Abiri, 2024). Statutory and constitutional interpretation methods (Scalia and Garner, 2012) are also important and have been adapted as tools for resolving ambiguity in specifications (Kolt et al., 2026).

Taxonomies of Alignment Approaches. Existing surveys of alignment and governance work generally organize the field by technical methods such as RLHF (Lu et al., 2025; Shen et al., 2023), or by value-setting principles (Chalkidis, 2025). For example, one line of work distinguishes between Overton, steerable, and distributional pluralism (Sorensen et al., 2024). Others have organized the field into pluralistic, contextual, and collective mechanisms (Bao et al., 2026). Still, the field remains largely unmapped, with the various approaches emerging largely independently of each other and not in direct conversation with one another.

Typology Construction. Motivated by taxonomies sorted by non-equivalent criteria and by free-floating typologies whose dimensions are never made explicit, prior work provides a template for building typologies from an overarching concept into explicit, clearly defined dimensions of an orthogonal nature (Collier et al., 2012). The underlying notion of a property space is due to earlier work in the same tradition (Barton, 1955),

as is the distinction between conceptual and explanatory typologies (Elman, 2005). As we find equivalent failure types in existing taxonomies of alignment, we follow this methodology.

Comparison to Related Work. Existing work maps traditional concepts from philosophical, political, and legal theory onto alignment and governance methods one at a time. We take a different approach by organizing the field according to a more complete set of underlying principles by which these methods organize themselves, and map individual traditions and concepts onto them subsequently. Furthermore, we attempt to critically interrogate whether the stated inspiration in traditional work results in a faithful technical implementation, and describe ways in which gaps or confusion emerge. Additionally, this approach allows us to identify gaps in approaches to alignment and governance which remain largely unexplored in the technical literature.

3. A Normative Design Space for Alignment: A Taxonomy

[roma: this needs to be the biggest section. we need to clarify how the taxonomy was created/why this current form is the best]

In order to study the use of normative concepts across fields, we create a two-dimensional property space defined by two axes: the locus of normative authority and the disposition of normative conflict.

We derive these dimensions through a systematic mapping of alignment and governance literature followed by conceptual substruction: overlapping surface categories are decomposed into the underlying attributes that distinguish them. Following Petersen et al. (2008), two coders independently extract normative concepts from the sampled literature and reconcile near-equivalent categories; full collection and inclusion criteria appear in Appendix A. Unlike conventional typologies that discard unrealized combinations, we retain empty cells because absence from the technical literature is itself analytically informative for identifying underexplored architectures.

We therefore follow prior work in using substruction (Barton, 1955) to recover underlying dimensions and extend them into a full property space here (Lazarsfeld, 1937). Substruction is "the procedure of finding, for a given system of types, the attribute space in which it belongs and the reduction which has been implicitly used" (Lazarsfeld, 1937). It is appropriate here because the categories which we previously derived overlap as projections of a higher-dimensional space onto a single axis. We deem this procedure applicable since prior work also considers normative types and arrives at axes similar to ours (Barton, 1955), as shown later.

This process yields two axes: locus of normative authority, which indicates where the binding normative content resides prior to the moment of decision, which here corresponds to model inference, and disposition of normative conflict, which indicates how any conflicts should be resolved. Their cross product gives a 5 x 3 property space. We note a difference from the null resolution procedure of prior work (Lazarsfeld, 1937). Whereas that procedure removes any null cells, we retain them, since we are interested in gaps in the literature.

For the locus of authority axis, we find the following categories:

1. **Codified.** The norms the AI ought to follow exist as an explicit written specification, set prior to deployment and, in principle, complete before any particular case arises. Authority resides in this specification (Hart, 2012).
2. **Case-accumulated.** The norms reside in a growing body of past decided cases. No case is decided by a rule stated in advance; what the AI should do is whatever the accumulated decisions have held (Jonsen and Toulmin, 1988; Levi, 2013).
3. **Customary.** The norms reside in regularities of practice (what agents in the population actually do), never written down and never adjudicated. There is no authoritative moment at which a norm gets explicitly stated (Burke, 1986; Lewis, 2008).
4. **Dispositional.** The norms reside in the AI's

trained character rather than in any external artifact. There is no rule to consult and nothing to look up, only the disposition itself (Aristotle).

5. **Procedural.** No norms persist between decisions. What the AI should do is whatever a specified procedure (e.g., a vote, a deliberation, a debate) yields at the moment of decision (Habermas, 2015).

For the disposition of normative conflict, we find the following categories:

1. **Collapse.** When the AI's norms conflict, the system resolves the conflict into a single output, without trace or acknowledgment of the disagreement that produced it. The plurality is thereby aggregated, averaged, or adjudicated away (Sorensen et al., 2024).
2. **Contest.** The conflict between norms is preserved and kept procedurally live. Either the output itself carries the disagreement, or the decision stays open to challenge through a standing channel (Mouffe, 2000; Pettit, 1997).
3. **Partition.** The conflict is neither resolved nor sustained but routed around. Multiple normative orders are kept intact, and a rule assigns which one governs a given case (de Sousa Santos, 2002; Lijphart, 1977; Savigny, 1867).

We considered open, axial and selective coding (Corbin and Strauss, 2014), as well as the first-order, second-order and aggregate-dimension ladder (Gioia et al., 2013), which would perhaps read as more familiar to a social science reader. We prefer the keyword-based approach with dimension resolution, as it is more native to the computer science literature it addresses. However, we note that the alternatives would likely yield a comparable outcome in terms of the intermediate categories.

General morphological analysis (Ritchey, 2022) would be inappropriate here, given that it constructs its multi-dimensional space while inspecting unrealized combinations, which requires preemptive stipulation. Since we are working with

an existing structure and literature landscape, this would introduce unnecessary assumptions and likely lead to a less grounded result.

4. Case Studies in the Design Space

[roma: this can cover 2-3 case studies in about 1 page.]

We use our 53 typology to examine a few different alignment methods that are widely used in the literature. Specifically, we ask the question: **how have normative concepts been translated into technical systems and what assumptions are made?** Our purpose is not just to evaluate the fidelity of technical implementations to their philosophical sources, but to examine what functions of the original normative framework are preserved, as well as how design choices affect the resulting governance structure.

4.1. Preference-based Alignment

Preference-based reward feedback to align LLMs is a widely used paradigm for AI systems. A standard approach involves training a model to infer human preferences over LLM outputs, and then using this model within a reinforcement learning paradigm to steer LLM outputs towards the approximated aggregate preferences. The resulting model is meant to be *aligned* to users as a result of this optimization towards the aggregate. This idea has roots in utilitarian moral philosophy—Bentham's principle of utility Bentham (1994) evaluates decisions according to their effect on the aggregate happiness of humans affects, while Mill posits that individual welfare contributes to an overall social good Mill (2011).

The many critiques against this include Tocqueville and Mill that warn against the *tyranny of the majority* i.e., a resulting state where aggregate or majoritarian consensus can erase or suppress minority perspectives and contradicting world views. Later work in social choice theory and welfare economics provide theoretical arguments against the validity of these frameworks showing how no rank-order voting systems—such as those used in preference algorithms—can convert individual preferences into a population-level consensus while maintaining fairness axioms ?. Al-

though preference-based alignment aims to elicit comparative judgements from humans at data collection time, they ultimately contribute to a common reward function, rank ordering, or model policy—thus only representing the aggregate population vote. Technical translations therefore make an architectural choice to represent plurality upstream but then collapse it downstream. In our taxonomy, this aggregation falls directly under *collapse*: where competing judgements are transformed into a common ordering. **[roma: spend at least 2-3 sentences talking about disagreement work, but how it hasn't made it into mainstream alignment techniques]** There do exist alternative architectures in which disagreement survives, and where minority judgements or conflicting normative claims can remain represented downstream. There is a big design opportunity therefore for systems that treat disagreement as information to gain signal from rather than noise to be eliminated.

4.2. Constitutions and Rules

Moving away from utilitarian preference aggregation, CAI turns to statutory constitutionalism and legal positivism, which sees law as a system of explicit rules created by a recognized authority, whose legitimacy stems from the formal correctness and enactment of said rules, not necessarily a consensus from involved parties. In the terminology of our taxonomy, they therefore locate normative authority primarily in a codified artifact: the constitution exists prior to a particular model decision and supplies normative guidance for evaluating it.

Yet a constitution in this technical sense performs only some of the functions associated with constitutions in legal and political theory. Constitutional governance ordinarily concerns not only the existence of a written text, but also questions about who may interpret it, how conflicts among provisions are resolved, how its application develops through precedent, under what conditions it may be amended, and which institutions possess authority to contest or revise decisions. A fixed specification largely separates the normative text from this surrounding institutional machinery.

This translation is not necessarily a defect. Reducing a constitution to an explicit specification makes normative guidance legible, auditable, and usable within existing training pipelines. But it also changes the normative architecture. Authority that is distributed across text, interpretation, precedent, and institutional roles becomes concentrated in the specification itself and in the process that authors it. Disagreement that constitutional institutions may keep procedurally live can instead be resolved before deployment, when principles are selected and encoded.

This comparison therefore suggests a different question from whether Constitutional AI is “sufficiently constitutional”: which functions of constitutional governance are useful for alignment? Existing work has emphasized the specification of principles. Less explored possibilities include mechanisms for authoritative interpretation, amendment, appeal, precedent, or explicit resolution of conflicts among constitutional principles. Normative translation makes these omissions visible as design choices rather than treating “a constitution” as a unitary technical object.

4.3. Deliberation and Collective Judgements

Deliberative approaches provide a third form of translation. Rather than locating normative authority in a fixed text or directly aggregating independently elicited preferences, these methods generate judgments through a procedure in which participants exchange reasons, revise positions, or construct principles collectively. They therefore move normative authority toward the procedural end of our first dimension. Here again, however, the technical mechanism can preserve one component of the source concept while transforming others. In theories of deliberative democracy, deliberation is not simply communication among multiple agents. Its normative force depends on properties of the procedure: who participates, whose reasons receive consideration, what forms of justification are admissible, whether participants can challenge decisions, and under what conditions an outcome is considered legitimate. A computational system in which humans or models exchange arguments may reproduce the surface mechanism of deliberation without necessarily

reproducing these conditions.

More importantly for our second dimension, deliberation does not uniquely determine what happens to disagreement. A deliberative procedure may seek convergence on a single principle or policy, in which case it remains a collapse architecture. Alternatively, it may preserve dissent, minority reports, or opportunities for subsequent challenge, moving toward contest. Thus, labeling a system “deliberative” is insufficient to characterize its normative architecture: the same procedural mechanism can instantiate different treatments of normative conflict. This illustrates why separating our two dimensions is useful. Where normative authority comes from does not determine what a system must do with disagreement. Procedural systems need not collapse, just as codified systems need not be uncontestable. Normative theory makes these combinations visible even where current technical implementations concentrate on only a subset of them.

4.4. Opposition and Adversaries in Red-teaming

A different approach to normative governance begins not from agreement, but from institutionalized opposition. Adversarial legal procedures assign distinct actors the role of challenging claims, evidence, and interpretations before a decision is reached; political theories of contestability similarly emphasize the ability of affected parties to challenge exercises of authority rather than requiring consensus over every decision. In both cases, disagreement is not merely an obstacle to good decision-making: structured opposition can itself provide information that would otherwise remain hidden.

Adversarial red-teaming instantiates a structurally similar idea in AI systems. Human or automated adversaries actively search for behaviors that violate safety objectives, expose inconsistencies, or reveal failures that ordinary evaluation may not surface. Recent multi-agent approaches extend this structure by training or prompting models to generate increasingly effective attacks, critiques, or counterarguments against other models. Unlike preference aggregation, which seeks

to combine conflicting judgments into a common objective, adversarial methods derive their value precisely from maintaining opposing roles during evaluation. However, most red-teaming systems use contestation primarily as an epistemic instrument: opposition is valuable because it discovers failures, after which those failures can be incorporated into training, filtering, or evaluation. The contest itself generally does not remain part of the system’s normative decision procedure. Once an attack exposes undesirable behavior, the resulting information is typically folded back into a revised model or policy, returning the architecture toward collapse.

Normative theories of adversarial procedure and contestability suggest a stronger alternative. Contestation can be understood not only as a mechanism for discovering errors before deployment, but as a persistent feature of legitimate decision-making. An alignment architecture could preserve adversarial roles at inference time, allow decisions or normative justifications to be challenged, record successful challenges, and specify procedures through which contested decisions are revised. Such a system would move from red-teaming the model toward making the model’s normative decisions themselves contestable. This distinction reveals an underexplored transition in the design space: alignment already has sophisticated machinery for generating adversarial disagreement, but primarily uses it to improve a subsequently unified policy. The same machinery could instead support architectures in which disagreement remains procedurally live. Red-teaming therefore provides a natural technical bridge from existing alignment practice toward contest-oriented normative governance.

5. Expanding Design Space

[roma: this can discuss what is missing and future gaps to fill] Across these examples, technical translation tends to preserve the component of a normative institution that can most readily be represented as a computational mechanism: constitutions become specifications, heterogeneous preferences become aggregation problems, and delib-

eration becomes a procedure for producing judgments. The surrounding mechanisms through which normative institutions manage interpretation, disagreement, revision, and competing authority are less consistently retained.

We do not interpret this pattern as evidence that existing alignment methods are incorrect implementations of their theoretical analogues. Computational systems face different constraints and objectives from legal or political institutions, and simplification is often necessary. Rather, the comparison exposes where normative decisions have moved. A fixed constitution moves questions of interpretation and conflict toward specification design; an aggregate preference model moves disagreement into the aggregation rule; and consensus-oriented deliberation moves conflict into the procedure that determines its outcome.

This observation also connects normative translation to the gaps in our property space. Components omitted during translation frequently resemble the mechanisms that distinguish contest and partition architectures: appeal, amendment, persistent dissent, jurisdiction, competing sources of authority, and procedures for revisiting prior decisions. The sparsity of these regions may therefore reflect not an absence of normative resources for designing such systems, but the selective translation of those resources into alignment to date.

The value of comparison with normative theory is consequently generative rather than merely diagnostic. It allows us to ask not whether alignment systems faithfully reproduce existing institutions, but which institutional mechanisms become available once we stop treating their current technical translations as exhaustive. We turn next to these underexplored architectures and the design opportunities they suggest

Our analysis reveals two sources of underexplored alignment designs. Architectural gaps arise where few or no existing methods occupy a region of the property space, suggesting combinations of normative authority and conflict management that have received limited technical attention. Mechanism gaps arise within existing architectures when potentially useful components

of their normative analogues—such as contestation, precedent, amendment, or jurisdiction—are absent from technical implementations. Neither type of gap implies that the missing design should be adopted: some may be undesirable or impractical for AI systems. Rather, they provide structured hypotheses about alternative alignment mechanisms whose tradeoffs can be implemented and evaluated.

6. Conclusions

[roma: also either bleed in limitations here, or have a section for it before this]

AI alignment is not only a problem of identifying which values a model should represent, but also of deciding how normative authority should be structured when those values conflict. We introduced a taxonomy that makes these choices explicit, organizing alignment approaches by where normative authority resides and how normative conflict is treated. Viewing existing systems through this lens also makes visible the consequences of *normative translation*: technical implementations preserve some features of the philosophical, political, and legal traditions they draw from while leaving others behind. Making these design choices explicit expands the space of possible alignment approaches, particularly toward systems that can preserve, contest, or partition persistent disagreement rather than requiring it to collapse into a single normative objective.

References

- G. Abiri. Public constitutional AI. *Peking University School of Transnational Law Research Paper, Georgia Law Review*, 59, 2024.
- G. E. Anscombe. Modern moral philosophy. In *The definition of morality*, pages 211–234. Routledge, 2020.
- Anthropic. Claude’s constitution, a. URL <https://www.anthropic.com/constitution>.
- Anthropic. Claude’s new constitution, b. URL <https://www.anthropic.com/news/claude-new-constitution>.

- Aristotle. *Nicomachean Ethics*, volume 9 of *The Works of Aristotle Translated into English*. Clarendon Press.
- K. J. Arrow. *Social choice and individual values*, volume 12. Yale university press, 2012.
- Y. Bai et al. Constitutional AI: Harmlessness from AI feedback, 2022.
- H. Bao, Y. Huang, X. Wang, Z. Zhang, Y. Zhou, C. Yang, X. Zhang, and Y. Ye. AI alignment breaks at the edge. *arXiv preprint arXiv:2602.20042*, 2026.
- A. H. Barton. The concept of property-space in social research. *The language of social research*, 111:40, 1955.
- J. Bentham. Of the principle of utility. *Ethics*, 20 (4):306–312, 1994.
- I. Berlin. *Four Essays on Liberty*. Oxford University Press, London and New York, 1969.
- I. Berlin. Two concepts of liberty. In *Reading political philosophy*, pages 231–237. Routledge, 2014.
- E. Burke. Reflections on the revolution in france (1790). *Everyman’s Library*, 1986.
- N. A. Caputo. Alignment as jurisprudence. *Yale Journal of Law and Technology (forthcoming)*, 2024.
- I. Chalkidis. Decoding alignment: A critical survey of LLM development initiatives through value-setting and data-centric lens. *arXiv preprint arXiv:2508.16982*, 2025.
- J. Cohen. Deliberation and democratic legitimacy. In *Debates in contemporary political philosophy*, pages 352–370. Routledge, 2005.
- D. Collier, J. LaPorte, and J. Seawright. Putting typologies to work: Concept formation, measurement, and analytic rigor. *Political research quarterly*, 65(1):217–232, 2012.
- J. Corbin and A. Strauss. *Basics of qualitative research: Techniques and procedures for developing grounded theory*. Sage publications, 2014.
- B. de Sousa Santos. Toward a new legal common sense: Law. *Globalization and Emancipation*, 314, 2002.
- R. Dworkin. *Law’s empire*. Harvard university press, 1988.
- C. Elman. Explanatory typologies in qualitative studies of international politics. *International Organization*, 59(2):293–326, 2005. doi: 10.1017/S0020818305050101.
- S. Feng, T. Sorensen, Y. Liu, J. Fisher, C. Y. Park, Y. Choi, and Y. Tsvetkov. Modular pluralism: Pluralistic alignment via multi-LLM collaboration. In Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 4151–4171, Miami, Florida, USA, Nov. 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.240. URL <https://aclanthology.org/2024.emnlp-main.240/>.
- L. Ge, D. Halpern, E. Micha, A. D. Procaccia, I. Shapira, Y. Vorobeychik, and J. Wu. Axioms for ai alignment from human feedback. *Advances in Neural Information Processing Systems*, 37:80439–80465, 2024.
- D. A. Gioia, K. G. Corley, and A. L. Hamilton. Seeking qualitative rigor in inductive research: Notes on the gioia methodology. *Organizational research methods*, 16(1):15–31, 2013.
- J. Habermas. *Between Facts and Norms: Contributions to a Discourse Theory of Law and Democracy*. MIT Press, Cambridge, MA, 1996.
- J. Habermas. *Between facts and norms: Contributions to a discourse theory of law and democracy*. John Wiley & Sons, 2015.
- H. L. A. Hart. *The Concept of Law*. Clarendon Press, Oxford, 1961.
- H. L. A. Hart. *The concept of law*. OUP Oxford, 2012.
- T. Hobbes. *Leviathan, or, The matter, forme, & power of a common-wealth ecclesiasticall and civil*, volume 21. 1885.

- S. Huang, D. Siddarth, L. Lovitt, T. I. Liao, E. Durmus, A. Tamkin, and D. Ganguli. Collective constitutional AI: Aligning a language model with public input. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1395–1417, 2024.
- A. R. Jonsen and S. Toulmin. *The abuse of casuistry: A history of moral reasoning*. Univ of California Press, 1988.
- N. Kolt, N. Caputo, J. Boeglin, C. O’Keefe, R. Bommasani, S. Casper, M.-F. Cuéllar, N. Feldman, I. Gabriel, G. K. Hadfield, et al. Legal alignment for safe and ethical AI. *arXiv preprint arXiv:2601.04175*, 2026.
- P. F. Lazarsfeld. Some remarks on the typological procedures in social research. *Zeitschrift für Sozialforschung*, 6(1):119–139, 1937.
- E. H. Levi. *An introduction to legal reasoning*. University of Chicago Press, 2013.
- D. Lewis. *Convention: A philosophical study*. John Wiley & Sons, 2008.
- A. Lijphart. *Democracy in plural societies: A comparative exploration*. Yale University Press, 1977.
- K. N. Llewellyn. *The common law tradition: deciding appeals*, volume 16. Quid Pro Books, 2016.
- J. Locke. *Two Treatises of Government*. Cambridge Texts in the History of Political Thought. Cambridge University Press, student edition.
- R. Long, J. Sebo, P. Butlin, K. Finlinson, K. Fish, J. Harding, J. Pfau, T. Sims, J. Birch, and D. Chalmers. Taking AI welfare seriously. *arXiv preprint arXiv:2411.00986*, 2024.
- H. Lu, L. Fang, R. Zhang, X. Li, J. Cai, H. Cheng, L. Tang, Z. Liu, Z. Sun, T. Wang, et al. Alignment and safety in large language models: Safety mechanisms, training paradigms, and emerging challenges. *arXiv preprint arXiv:2507.19672*, 2025.
- R.-R. Maura-Rivero, C. Nagpal, R. Patel, and F. Visin. Utility-inspired reward transformations improve reinforcement learning training of language models, 2025. URL <https://arxiv.org/abs/2501.06248>.
- J. S. Mill. *Utilitarianism ()*. Hyweb Technology Co. Ltd., 2011.
- C. Mouffe. *The democratic paradox*. verso, 2000.
- A. Ovadya, K. Redman, L. Thorburn, Q. Z. Chen, O. Smith, F. Devine, A. Konya, S. Milli, M. Revel, K. Feng, et al. Democratic AI is possible. the democracy levels framework shows how it might work. *arXiv preprint arXiv:2411.09222*, 2024.
- D. Parfit. *Reasons and persons*. Oxford University Press, 1987.
- K. Petersen, R. Feldt, S. Mujtaba, and M. Mattsson. Systematic mapping studies in software engineering. In *12th international conference on evaluation and assessment in software engineering (EASE)*. BCS Learning & Development, 2008.
- P. Pettit. *Republicanism: A theory of freedom and government*. Oxford University Press, 1997.
- I. Rahwan. Society-in-the-loop: programming the algorithmic social contract. *Ethics and information technology*, 20(1):5–14, 2018.
- J. Rawls. *Political Liberalism*. Columbia University Press, New York, 1993.
- J. Rawls. A theory of justice. In *Applied ethics*, pages 21–29. Routledge, 2017.
- T. Ritchey. General morphological analysis: an overview. 2022.
- W. D. Ross. *The Right and the Good*. Clarendon Press, Oxford, 1930.
- J.-J. Rousseau. *The Social Contract and Other Later Political Writings*. Cambridge Texts in the History of Political Thought. Cambridge University Press.
- F. K. v. Savigny. System of the modern roman law. (No Title), 1867.

A. Scalia and B. A. Garner. *Reading Law: The Interpretation of Legal Texts*. Thomson/West, St. Paul, MN, 2012. ISBN 978-0-314-27555-4.

A. Sen. *Collective choice and social welfare: Expanded edition*. Penguin UK, 2017.

T. Shen, R. Jin, Y. Huang, C. Liu, W. Dong, Z. Guo, X. Wu, Y. Liu, and D. Xiong. Large language model alignment: A survey. *arXiv preprint arXiv:2309.15025*, 2023.

N. Soares, B. Fallenstein, E. Yudkowsky, and S. Armstrong. Corrigibility. In *AAAI Workshops: Workshops at the Twenty-Ninth AAAI Conference on Artificial Intelligence*, Austin, TX, 2015. AAAI Publications.

T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Miresghallah, C. M. Rytting, A. Ye, L. Jiang, X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024.

T. Sorensen, P. Mishra, R. Patel, M. H. Tessler, M. A. Bakker, G. Evans, I. Gabriel, N. Goodman, and V. Rieser. Value profiles for encoding human variation. In C. Christodoulopoulos, T. Chakraborty, C. Rose, and V. Peng, editors, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2047–2095, Suzhou, China, Nov. 2025. Association for Computational Linguistics. ISBN 979-8-89176-332-6. doi: 10.18653/v1/2025.emnlp-main.106. URL <https://aclanthology.org/2025.emnlp-main.106/>.

C. R. Sunstein. *Legal reasoning and political conflict*. Oxford University Press, 2018.

M. A. Warren. *Moral status: Obligations to persons and other living things*. Clarendon Press, 1997.

L. Weidinger, K. R. McKee, R. Everett, S. Huang, T. O. Zhu, M. J. Chadwick, C. Summerfield, and I. Gabriel. Using the veil of ignorance to align AI systems with principles of justice. *Proceedings of the National Academy of Sciences*, 120(18): e2213709120, 2023.

Acknowledgements

7. Appendix