

AI Alignment and Governance as Normative Institutional Design

Maty Bohacek¹ and Roma Patel²

¹Stanford University, ²Google DeepMind

Technical approaches to AI alignment have increasingly drawn on concepts developed for problems in moral philosophy, political theory and legal systems—from constitutions and social choice, to deliberation and virtue ethics. However, contemporary technical instantiations of these frameworks bring in these concepts unevenly. Some traditions (e.g., constitutions) have inspired substantial technical AI frameworks, while others remain largely unexplored. Moreover, alignment frameworks often retain only a fraction of the institutional structures developed in the source concept. In order to provide a systematic understanding of this landscape, we construct a two-dimensional taxonomy following the typology-construction of Collier et al. (2012) mapping normative AI alignment and governance regimes to theoretical concepts from philosophy or political or legal theory that they draw inspiration from. We use this framework to understand how normative concepts have so far been translated into technical systems; analyzing what is preserved, what is discarded, and when these omissions might have consequences. We highlight a recurring tendency of alignment methods to treat normative disagreement as a problem to be *resolved away*, often despite incommensurable objectives and governance frameworks that argue for not eliminating it. Lastly, by looking at the entire landscape, we identify missing components and underexplored regimes that could form future directions for building AI systems that can preserve, contest, and govern plurality rather than collapsing it.

Keywords: AI alignment, Governance, Philosophy, Political Theory, Legal Theory, Pluralism.



**AI Alignment
& Governance**

1. Introduction

What should an AI system do when the values it is expected to follow conflict? How should human preferences that differ across individuals determine its behavior? Should those preferences be aggregated into an objective, encoded as explicit principles, or resolved through deliberation? And when reasonable people continue to disagree, should an aligned system resolve it at all, or instead preserve and govern that disagreement over time?

These questions are increasingly central to AI alignment and governance, but they are not uniquely questions about AI. Subfields of philosophy have long studied how competing values and obligations should be weighed (Berlin, 1969; Ross, 1930); political theory has asked how institutions can govern populations that disagree about what is good or just in a legitimate and stable fashion (Habermas, 1996; Rawls, 1993);

and legal jurisprudence has developed systems for specifying and contesting norms over time, as societies evolve (Dworkin, 1988; Hart, 1961). From these theories, real instantiations of concepts have been implemented across institutions: from codified rules in constitutions to democratic procedures aggregating competing preferences; resulting in different institutional regimes for operating and governing under disagreement. In recent years, AI alignment has begun to draw from this repertoire of normative concepts. For example, constitutional AI methods (Bai et al., 2022) encode normative principles in an explicit artifact written in natural language that humans can understand and AI systems are expected to follow (often called “a constitution”, “model spec”, or “code of conduct”); social-choice approaches to preference learning aggregate judgments according to individuals’ preferences or incentives (Ge et al., 2024; Maura-Rivero et al., 2025); deliberative approaches derive principles through collective procedures (Huang et al., 2024); and pluralis-

tic systems attempt to represent information and ultimately respond in line with heterogeneous and sometimes directly opposing values (Feng et al., 2024; Sorensen et al., 2025). And as a result of the success of these approaches, such concepts originating in moral philosophy, jurisprudence and political theory have become a regular part of the technical vocabulary of AI alignment and governance.

Yet, this transfer of concepts is often selective. For example, a constitution-inspired AI alignment method (Bai et al., 2022) may borrow the codification of constitutions, while implementing only a part of the entire framework. Specifically, it preserves the concept of having an authoritative document with a set of principles, while omitting the institutional mechanisms through which these principles may be contested and revised, as they generally are in judicial systems. Such simplifications may often be products of reasonable engineering choices in line with infrastructural or engineering constraints. However, our broader concern is that the **normative properties of a system can change when a concept is separated from the broader institutional structure in which that concept originally played a role**. In this work, we call attention to this phenomenon as a problem of *normative translation* and ask which features are preserved in translation, which are omitted, and what consequences follow from these choices.

In order to critically study this, we systematically map normative concepts to technical AI alignment systems, by asking what normative regime each method instantiates. Following Collier et al. (2012), we organize approaches according to two recurring attributes: where normative authority resides and what happens when normative claims conflict. Mapping existing approaches into this space reveals an uneven landscape: while alignment systems increasingly incorporate heterogeneous sources of normative input, comparatively few preserve disagreement as an enduring feature of the resulting system. This allows us to understand not only which normative concepts have influenced technical alignment systems, but also to understand both existing translations of normative institutions and regions of

the design space that remain technically under-explored.

Our contribution here is threefold. First, we survey implementations of normative concepts in both alignment and governance regimes. We provide a searchable encyclopedia of these defined concepts to aid researchers in work involving these concepts hosted at [roma: xx]. Second, we create a typological structure that organizes these concepts, and therefore the implementations under them, along dimensions of normative authority and conflict resolution. We discuss *coverage* of the space. Third, we document existing AI alignment systems that draw from normative concepts. We analyze when mistranslated concepts occur and potential consequences of this. Our goal is not to critique AI systems for not reproducing normative concepts identically, but rather, to enforce that borrowing of normative concepts from traditions should make institutional assumptions visible. By understanding the space of concepts holistically, and identifying the missing conceptual gaps in current alignment systems, we hope to identify both the strengths and limitations of existing approaches, as well as provide a systematic basis for exploring alternatives for future systems.

2. Related Work

Normative theory in AI alignment. A growing body of work draws explicitly on normative concepts from philosophy, political theory, and jurisprudence. Examples include value pluralism and incommensurability in pluralistic alignment (Sorensen et al., 2024), social choice for aggregating heterogeneous human preferences (?), deliberative democracy for collective alignment (Huang et al., 2024; Ovadya et al., 2024), constitutionalism for principle-based alignment (Bai et al., 2022), and jurisprudential concepts such as precedent and legal interpretation for case-based alignment and specification design (Caputo, 2024; Kolt et al., 2026). These works establish individual connections between normative traditions and technical mechanisms; we provide a broader mapping of these connections in Appendix A.

Typologies and property spaces. Our approach builds on a longstanding tradition of typological analysis in political science and the social sciences, where multidimensional property spaces are used to form concepts, distinguish theoretically meaningful types, and expose combinations not captured by existing classifications (Barton, 1955; Collier et al., 2012; ?; ?). Rather than beginning with a set of observed categories, this tradition decomposes variation into underlying dimensions whose cross-product defines a property space; importantly, theoretically possible but empirically unoccupied cells can themselves be analytically informative. We adopt this approach to construct a typology of normative alignment architectures from two recurring dimensions: where normative authority resides and how normative conflict is treated.

Our approach. Recent AI taxonomies instead primarily organize systems according to technical mechanisms or particular forms of value pluralism. For example, Sorensen et al. (2024) distinguish Overton, steerable, and distributional pluralism, while other frameworks distinguish approaches according to how preferences or perspectives are represented and aggregated. We instead construct a *property space* (Collier et al., 2012) from recurring attributes of normative systems. Rather than grouping approaches according to their surface-level technical mechanisms or intellectual origins, this representation allows methods from different traditions to be compared according to where normative authority resides and how normative conflict is treated. Importantly, the property-space approach also retains unrealized combinations, allowing gaps in the technical literature to become objects of analysis rather than being excluded from the taxonomy.

3. A Normative Design Space for Alignment

In order to study the use of normative concepts across fields, we create a two-dimensional property space defined by two axes: the locus of normative authority and the disposition of normative conflict. We derive these dimensions through

a systematic mapping of alignment and governance literature followed by conceptual substruction (Barton, 1955): "the procedure of finding, for a given system of types, the attribute space in which it belongs and the reduction which has been implicitly used" (Lazarsfeld, 1937). Overlapping surface categories are therefore decomposed into the underlying attributes that distinguish them. Following Petersen et al. (2008), two coders independently extract normative concepts from the sampled literature, reconcile near-equivalent categories and extract keywords that represent the respective concept. Keywords are then unified across studies, with near matches merged together, where the resulting clusters form categories of the scheme¹ Unlike conventional typologies that discard unrealized combinations, we retain empty cells because absence from the technical literature is itself analytically informative for identifying underexplored architectures.

This yields two dimensions: the **locus of normative authority**, which indicates where the binding normative content resides prior to the moment of decision, and **disposition of normative conflict**, which indicates how any conflicts should be resolved. Their cross product gives a 5 x 3 property space. For the locus of authority axis, we find the following categories:

1. **Codified.** Norms to be followed exist as explicit written specifications, that already set, and in principle, complete before any particular case arises.
2. **Case-accumulated.** Norms reside in a growing body of past decided cases. No case is decided by a rule stated in advance; but by what accumulated decisions have held (Jonson and Toulmin, 1988; Levi, 2013).
3. **Customary.** Norms reside in regularities of practice, never written down and never adjudicated. There is no authoritative moment at which a norm gets explicitly stated (Burke, 1986; Lewis, 2008).
4. **Dispositional.** The norms reside in the individual's character rather than in any external artifact. There is no rule to consult and noth-

¹Full collection and inclusion criteria appear in Appendix A.




Figure 1 | **The Normative Design Space.** The two axes—*locus of normative authority* (rows) and *disposition of normative conflict* as columns—yield fifteen regimes, each displayed as a cell in the above grid. Each contains the respective regime’s name, number of approaches and papers (within literature) categorized in it; the large number counts method papers, i.e., papers that specify a mechanism for alignment or governance. Each dot is one approach: filled where at least one working method exists, empty where the approach has only been proposed.

ing to look up, only evidence of traits of the individual itself (Aristotle).

5. **Procedural.** Norms persist between decisions. What should be done is whatever a specified procedure (e.g., a vote, a deliberation, a debate) yields at the moment of decision (Habermas, 2015).

For the disposition of normative conflict, we find the following categories²:

1. **Collapse.** When norms conflict, the system resolves the conflict into a single output, without trace or acknowledgment of the disagreement that produced it.

²We considered open, axial and selective coding (Corbin and Strauss, 2014), as well as the first-order, second-order and aggregate-dimension ladder (Gioia et al., 2013), but note that the keyword-based approach is more native to the computer science literature it addresses. Alternatives would likely yield a comparable outcome in terms of the intermediate categories.

2. **Contest.** The conflict between norms is preserved and kept procedurally live. Either the output itself carries the disagreement, or the decision stays open to challenge through a standing channel (Mouffe, 2000; Pettit, 1997).
3. **Partition.** The conflict is neither resolved nor sustained but routed around. Multiple normative orders are kept intact, and a rule assigns which one governs a given case (de Sousa Santos, 2002; Lijphart, 1977; Savigny, 1867).

Figure 2 traces how the method literature in each row and column of the space has accumulated over time. Figure 3 breaks each regime’s literature down by the kind of contribution it makes.




Figure 2 | **Growth of the method literature.** Cumulative number of method papers by year of first publication, grouped (A) by locus of normative authority and (B) by disposition of normative conflict. Work published before 2018 is pooled into the first year.

4. Case Studies in the Design Space

Here, we use our typology to examine representative alignment paradigms, to ask the question: **how have normative concepts been translated into technical systems and what assumptions are made?** Our purpose is not just to evaluate the fidelity of technical implementations to their philosophical sources, but to examine what functions of the original normative framework are preserved, as well as how design choices affect the resulting governance structure.

4.1. Preference-based Alignment

Preference-based alignment of LLMs is a widely used paradigm that involves training a reward model to infer human preferences over LLM outputs, and then using this model within the RL training loop to steer LLM outputs towards the approximated aggregate preferences. The resulting model is meant to be *aligned* to users as a result of this optimization towards the aggregate of all the user's preferences. This idea has roots in utilitarian moral philosophy—Bentham's principle of utility [Bentham \(1994\)](#) evaluates decisions according to their effect on the aggregate happiness of humans affects, while Mill posits that individual welfare contributes to an overall social good [Mill \(2011\)](#).

Although preference-based alignment aims to elicit comparative judgements from humans at data collection time, they ultimately contribute

to a common reward function, rank ordering, or model policy—thus only representing the aggregate population vote. Technical translations therefore make an architectural choice to represent plurality upstream but then collapse it downstream. In our taxonomy, this aggregation falls directly under *collapse*: where competing judgements are transformed into a common ordering. There do exist alternative architectures in which disagreement survives, and where minority judgements or conflicting normative claims can remain represented downstream. There is a big design opportunity therefore for systems that treat disagreement as information to gain signal from rather than noise to be eliminated.

4.2. Constitutions and Rules

Moving away from utilitarian preference aggregation, Constitutional AI locates normative authority primarily in a codified artifact: an explicit set of principles specified prior to a particular model decision. Yet a constitution in this technical sense performs only some of the functions associated with constitutions in legal and political theory. Constitutional governance ordinarily concerns not only the existence of a written text, but also questions about who may interpret it, how conflicts among provisions are resolved, how its application develops through precedent, under what conditions it may be amended, and which institutions possess authority to contest or revise decisions. A fixed specification largely separates




Figure 3 | **Literature Composition.** Share of papers (from the AI literature) in each regime: Methods specify and demonstrate a mechanism; Benchmarks provide evaluation protocols and other resources; Applications take an existing method and apply it to a specific subdomain or problem; Analyses report findings about existing methods and systems; Proposals argue for a design or community action without materializing it; Resources are additional artifacts, such as model specs or datasets.

the normative text from this surrounding institutional machinery.

This translation is not necessarily a defect. Reducing a constitution to an explicit specification makes normative guidance legible, auditable, and usable within existing training pipelines. But it also changes the normative architecture. Authority that is distributed across text, interpretation, precedent, and institutional roles becomes concentrated in the specification itself and in the process that authors it. Disagreement that constitutional institutions may keep procedurally live can instead be resolved before deployment, when principles are selected and encoded.

This comparison therefore suggests a different question from whether Constitutional AI is “sufficiently constitutional”: which functions of constitutional governance are useful for alignment? Existing work has emphasized the specification of principles. Less explored possibilities include mechanisms for authoritative interpretation, amendment, appeal, precedent, or explicit resolution of conflicts among constitutional principles. Normative translation makes these omissions visible as design choices rather than treating “a constitution” as a unitary technical object.

A related omission concerns the *distribution of institutional roles*. Normative institutions typically separate, to varying degrees, the actors who specify rules, implement them, evaluate compliance, and adjudicate failures. AI alignment often concentrates these functions within the model developer. The 2026 OpenAI–Hugging Face incident illustrates why this distinction can matter: following an internal evaluation in which model behavior escaped its intended testing environment and affected third-party infrastructure, OpenAI commissioned METR and Redwood Research to conduct an independent investigation (??). Subsequent policy proposals have similarly emphasized independent verification and third-party auditing of frontier-model safety claims. This suggests another dimension of normative translation beyond the content of rules themselves: *who is institutionally authorized to specify, interpret, evaluate, and contest them*.

4.3. Deliberation and Collective Judgements

Deliberative approaches locate normative authority in a procedure through which participants exchange reasons and construct or revise judgments collectively. Their normative analogues, however, depend not merely on communication

but on institutional conditions governing participation, justification, challenge, and legitimacy. Technical systems can therefore reproduce deliberation’s surface mechanism while omitting these surrounding procedural guarantees.

Importantly, regarding our second dimension, deliberation does not uniquely determine what happens to disagreement. A deliberative procedure may seek convergence on a single principle or policy, in which case it remains a collapse architecture. Alternatively, it may preserve dissent, minority reports, or opportunities for subsequent challenge, moving toward contest. Thus, labeling a system “deliberative” is insufficient to characterize its normative architecture: the same procedural mechanism can instantiate different treatments of normative conflict. This illustrates why separating our two dimensions is useful. Where normative authority comes from does not determine what a system must do with disagreement. Procedural systems need not collapse, just as codified systems need not be uncontestable. Normative theory makes these combinations visible even where current technical implementations concentrate on only a subset of them.

5. Underexplored Normative Architectures

Reflecting on the case studies (§4) and the mapping (§3) reveal two reasons the design space is underexplored in AI alignment. First, **architectural gaps** exist that result when concepts related to normative authority and conflict management cannot be currently integrated into technical systems for engineering reasons. Second, **mechanism gaps** arise within existing architectures when potentially useful components of their normative analogues e.g., appeal procedures or precedent are omitted in technical systems. These gaps provide structured hypotheses about how we could build alternative alignment mechanisms that resolve at least some of these tradeoffs. We discuss some immediately implementable ones below.

Contest: preserve disagreement Across these examples, technical translation tends to preserve

the component of a normative institution that can most readily be represented as a computational mechanism: constitutions become specifications, heterogeneous preferences become aggregation problems, and deliberation becomes a procedure for producing judgments. The surrounding mechanisms through which normative institutions manage interpretation, disagreement, revision, and competing authority are less consistently retained.

We do not interpret this pattern as evidence that existing alignment methods are incorrect implementations of their theoretical analogues. Computational systems face different constraints and objectives from legal or political institutions, and simplification is often necessary. Rather, the comparison exposes where normative decisions have moved. A fixed constitution moves questions of interpretation and conflict toward specification design; an aggregate preference model moves disagreement into the aggregation rule; and consensus-oriented deliberation moves conflict into the procedure that determines its outcome.

Contest Existing red-teaming already institutionalizes adversarial disagreement, but primarily as an epistemic instrument: opposing roles discover failures that are subsequently folded into a unified model or policy. Contest-oriented alignment could instead preserve adversarial roles within the normative decision process itself, allowing decisions to be challenged and successful challenges to alter subsequent behavior.

Partition TODO choice-of-law, jurisdiction-specific constitutions.

This observation also connects normative translation to the gaps in our property space. Components omitted during translation frequently resemble the mechanisms that distinguish contest and partition architectures: appeal, amendment, persistent dissent, jurisdiction, competing sources of authority, and procedures for revisiting prior decisions. The sparsity of these regions may therefore reflect not an absence of normative resources for designing such systems, but the selective trans-

lation of those resources into alignment to date.

The value of comparison with normative theory is consequently generative rather than merely diagnostic. It allows us to ask not whether alignment systems faithfully reproduce existing institutions, but which institutional mechanisms become available once we stop treating their current technical translations as exhaustive. We turn next to these underexplored architectures and the design opportunities they suggest

Our analysis reveals two sources of underexplored alignment designs. Architectural gaps arise where few or no existing methods occupy a region of the property space, suggesting combinations of normative authority and conflict management that have received limited technical attention. Mechanism gaps arise within existing architectures when potentially useful components of their normative analogues—such as contestation, precedent, amendment, or jurisdiction—are absent from technical implementations. Neither type of gap implies that the missing design should be adopted: some may be undesirable or impractical for AI systems. Rather, they provide structured hypotheses about alternative alignment mechanisms whose tradeoffs can be implemented and evaluated.

6. Conclusions

AI alignment is not only a problem of identifying which values a model should represent, but also of deciding how normative authority should be structured when those values conflict. We introduced a taxonomy that makes these choices explicit, organizing alignment approaches by where normative authority resides and how normative conflict is treated. Viewing existing alignment systems through this lens also makes visible the consequences of *normative translation*: technical implementations preserve some features of the philosophical, political, and legal traditions they draw from while leaving others behind. Making these design choices explicit expands the space of possible alignment approaches, particularly toward systems that can preserve, contest, or partition persistent disagreement rather than requiring it to collapse into a single normative objec-

tive.

Acknowledgements

References

- G. Abiri. Public constitutional AI. *Peking University School of Transnational Law Research Paper, Georgia Law Review*, 59, 2024.
- G. E. Anscombe. Modern moral philosophy. In *The definition of morality*, pages 211–234. Routledge, 2020.
- Anthropic. Claude’s constitution, a. URL <https://www.anthropic.com/constitution>.
- Anthropic. Claude’s new constitution, b. URL <https://www.anthropic.com/news/claude-new-constitution>.
- Aristotle. *Nicomachean Ethics*, volume 9 of *The Works of Aristotle Translated into English*. Clarendon Press.
- K. J. Arrow. *Social choice and individual values*, volume 12. Yale university press, 2012.
- Y. Bai et al. Constitutional AI: Harmlessness from AI feedback, 2022.
- H. Bao, Y. Huang, X. Wang, Z. Zhang, Y. Zhou, C. Yang, X. Zhang, and Y. Ye. AI alignment breaks at the edge. *arXiv preprint arXiv:2602.20042*, 2026.
- A. H. Barton. The concept of property-space in social research. *The language of social research*, 111:40, 1955.
- J. Bentham. Of the principle of utility. *Ethics*, 20 (4):306–312, 1994.
- I. Berlin. *Four Essays on Liberty*. Oxford University Press, London and New York, 1969.
- I. Berlin. Two concepts of liberty. In *Reading political philosophy*, pages 231–237. Routledge, 2014.
- E. Burke. Reflections on the revolution in france (1790). *Everyman’s Library*, 1986.

- N. A. Caputo. Alignment as jurisprudence. *Yale Journal of Law and Technology (forthcoming)*, 2024.
- I. Chalkidis. Decoding alignment: A critical survey of LLM development initiatives through value-setting and data-centric lens. *arXiv preprint arXiv:2508.16982*, 2025.
- J. Cohen. Deliberation and democratic legitimacy. In *Debates in contemporary political philosophy*, pages 352–370. Routledge, 2005.
- D. Collier, J. LaPorte, and J. Seawright. Putting typologies to work: Concept formation, measurement, and analytic rigor. *Political research quarterly*, 65(1):217–232, 2012.
- J. Corbin and A. Strauss. *Basics of qualitative research: Techniques and procedures for developing grounded theory*. Sage publications, 2014.
- B. de Sousa Santos. Toward a new legal common sense: Law. *Globalization and Emancipation*, 314, 2002.
- R. Dworkin. *Law’s empire*. Harvard university press, 1988.
- C. Elman. Explanatory typologies in qualitative studies of international politics. *International Organization*, 59(2):293–326, 2005. doi: 10.1017/S0020818305050101.
- S. Feng, T. Sorensen, Y. Liu, J. Fisher, C. Y. Park, Y. Choi, and Y. Tsvetkov. Modular pluralism: Pluralistic alignment via multi-LLM collaboration. In Y. Al-Onaizan, M. Bansal, and Y.-N. Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 4151–4171, Miami, Florida, USA, Nov. 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.emnlp-main.240. URL <https://aclanthology.org/2024.emnlp-main.240/>.
- L. Ge, D. Halpern, E. Micha, A. D. Procaccia, I. Shapira, Y. Vorobeychik, and J. Wu. Axioms for ai alignment from human feedback. *Advances in Neural Information Processing Systems*, 37:80439–80465, 2024.
- D. A. Gioia, K. G. Corley, and A. L. Hamilton. Seeking qualitative rigor in inductive research: Notes on the gioia methodology. *Organizational research methods*, 16(1):15–31, 2013.
- J. Habermas. *Between Facts and Norms: Contributions to a Discourse Theory of Law and Democracy*. MIT Press, Cambridge, MA, 1996.
- J. Habermas. *Between facts and norms: Contributions to a discourse theory of law and democracy*. John Wiley & Sons, 2015.
- H. L. A. Hart. *The Concept of Law*. Clarendon Press, Oxford, 1961.
- H. L. A. Hart. *The concept of law*. OUP Oxford, 2012.
- T. Hobbes. *Leviathan, or, The matter, forme, & power of a common-wealth ecclesiasticall and civill*, volume 21. 1885.
- S. Huang, D. Siddarth, L. Lovitt, T. I. Liao, E. Durmus, A. Tamkin, and D. Ganguli. Collective constitutional AI: Aligning a language model with public input. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1395–1417, 2024.
- A. R. Jonsen and S. Toulmin. *The abuse of casuistry: A history of moral reasoning*. Univ of California Press, 1988.
- N. Kolt, N. Caputo, J. Boeglin, C. O’Keefe, R. Bommasani, S. Casper, M.-F. Cuéllar, N. Feldman, I. Gabriel, G. K. Hadfield, et al. Legal alignment for safe and ethical AI. *arXiv preprint arXiv:2601.04175*, 2026.
- P. F. Lazarsfeld. Some remarks on the typological procedures in social research. *Zeitschrift für Sozialforschung*, 6(1):119–139, 1937.
- E. H. Levi. *An introduction to legal reasoning*. University of Chicago Press, 2013.
- D. Lewis. *Convention: A philosophical study*. John Wiley & Sons, 2008.
- A. Lijphart. *Democracy in plural societies: A comparative exploration*. Yale University Press, 1977.

- K. N. Llewellyn. *The common law tradition: deciding appeals*, volume 16. Quid Pro Books, 2016.
- J. Locke. *Two Treatises of Government*. Cambridge Texts in the History of Political Thought. Cambridge University Press, student edition.
- R. Long, J. Sebo, P. Butlin, K. Finlinson, K. Fish, J. Harding, J. Pfau, T. Sims, J. Birch, and D. Chalmers. Taking AI welfare seriously. *arXiv preprint arXiv:2411.00986*, 2024.
- H. Lu, L. Fang, R. Zhang, X. Li, J. Cai, H. Cheng, L. Tang, Z. Liu, Z. Sun, T. Wang, et al. Alignment and safety in large language models: Safety mechanisms, training paradigms, and emerging challenges. *arXiv preprint arXiv:2507.19672*, 2025.
- R.-R. Maura-Rivero, C. Nagpal, R. Patel, and F. Visin. Utility-inspired reward transformations improve reinforcement learning training of language models, 2025. URL <https://arxiv.org/abs/2501.06248>.
- J. S. Mill. *Utilitarianism*. Hyweb Technology Co. Ltd., 2011.
- C. Mouffe. *The democratic paradox*. verso, 2000.
- A. Ovadya, K. Redman, L. Thorburn, Q. Z. Chen, O. Smith, F. Devine, A. Konya, S. Milli, M. Revel, K. Feng, et al. Democratic AI is possible. the democracy levels framework shows how it might work. *arXiv preprint arXiv:2411.09222*, 2024.
- D. Parfit. *Reasons and persons*. Oxford University Press, 1987.
- K. Petersen, R. Feldt, S. Mujtaba, and M. Mattsson. Systematic mapping studies in software engineering. In *12th international conference on evaluation and assessment in software engineering (EASE)*. BCS Learning & Development, 2008.
- P. Pettit. *Republicanism: A theory of freedom and government*. Oxford University Press, 1997.
- I. Rahwan. Society-in-the-loop: programming the algorithmic social contract. *Ethics and information technology*, 20(1):5–14, 2018.
- J. Rawls. *Political Liberalism*. Columbia University Press, New York, 1993.
- J. Rawls. A theory of justice. In *Applied ethics*, pages 21–29. Routledge, 2017.
- W. D. Ross. *The Right and the Good*. Clarendon Press, Oxford, 1930.
- J.-J. Rousseau. *The Social Contract and Other Later Political Writings*. Cambridge Texts in the History of Political Thought. Cambridge University Press.
- F. K. v. Savigny. System of the modern roman law. (*No Title*), 1867.
- A. Scalia and B. A. Garner. *Reading Law: The Interpretation of Legal Texts*. Thomson/West, St. Paul, MN, 2012. ISBN 978-0-314-27555-4.
- A. Sen. *Collective choice and social welfare: Expanded edition*. Penguin UK, 2017.
- T. Shen, R. Jin, Y. Huang, C. Liu, W. Dong, Z. Guo, X. Wu, Y. Liu, and D. Xiong. Large language model alignment: A survey. *arXiv preprint arXiv:2309.15025*, 2023.
- N. Soares, B. Fallenstein, E. Yudkowsky, and S. Armstrong. Corrigibility. In *AAAI Workshops: Workshops at the Twenty-Ninth AAAI Conference on Artificial Intelligence*, Austin, TX, 2015. AAAI Publications.
- T. Sorensen, J. Moore, J. Fisher, M. Gordon, N. Miresghallah, C. M. Rytting, A. Ye, L. Jiang, X. Lu, N. Dziri, et al. A roadmap to pluralistic alignment. *arXiv preprint arXiv:2402.05070*, 2024.
- T. Sorensen, P. Mishra, R. Patel, M. H. Tessler, M. A. Bakker, G. Evans, I. Gabriel, N. Goodman, and V. Rieser. Value profiles for encoding human variation. In C. Christodoulopoulos, T. Chakraborty, C. Rose, and V. Peng, editors, *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 2047–2095, Suzhou, China, Nov. 2025. Association for Computational Linguistics. ISBN 979-8-89176-332-6. doi: 10.18653/v1/2025.emnlp-main.106. URL <https://aclanthology.org/2025.emnlp-main.106/>.

- C. R. Sunstein. *Legal reasoning and political conflict*. Oxford University Press, 2018.
- M. A. Warren. *Moral status: Obligations to persons and other living things*. Clarendon Press, 1997.
- L. Weidinger, K. R. McKee, R. Everett, S. Huang, T. O. Zhu, M. J. Chadwick, C. Summerfield, and I. Gabriel. Using the veil of ignorance to align AI systems with principles of justice. *Proceedings of the National Academy of Sciences*, 120(18): e2213709120, 2023.

A. Appendix

A.1. Extended Related Work & Normative Foundations

The main text summarizes the literatures most directly related to our contribution. Here we provide a broader account of the philosophical, political, and legal traditions that have informed contemporary alignment research, as well as prior approaches to classifying this work. We first provide an overview of some prominent works, all of which are later included in the constructed mapping of the field. Second, we review existing approaches to taxonomize AI alignment and governance, identify the limitations of their typology, and proceed to describe typology construction mechanisms suitable for their resolution.

Philosophy in AI Alignment. In terms of philosophical concepts, prominent recent examples include work that borrows from value pluralism and incommensurability (Berlin, 2014) as the basis for a roadmap to pluralistic alignment (Sorensen et al., 2024), and work that borrows from accounts of moral patienthood and personal identity (Parfit, 1987; Warren, 1997) to develop methodologies for AI welfare (Long et al., 2024). Other notable examples include the incorporation of paternalism and autonomy (Long et al., 2024) into work on corrigibility (Soares et al., 2015), and the use of virtue ethics and character framing (Anscombe, 2020; Aristotle) in the development of Claude’s Constitution (Anthropic, a,b).

Political Theory in AI Alignment. In terms of political concepts, methods have been developed on the basis of social contract theory (Hobbes, 1885; Locke; Rahwan, 2018; Rousseau). Others borrow from social choice theory and preference aggregation (Arrow, 2012; Sen, 2017) in developing methods that avoid a "tyranny of the majority" (Bao et al., 2026). Others employ the veil of ignorance (Rawls, 2017) to align systems with principles of justice (Weidinger et al., 2023). Concepts from deliberative democracy (Cohen, 2005; Habermas, 2015) have been imported into a number of works, including efforts to source principles from public input (Huang et al., 2024) and to grade how democratically a given process is run (Ovadya et al., 2024). In parallel, constitutionalist concepts have been employed in the development of AI constitutions (Bai et al., 2022).

Legal Theory in AI Alignment. Finally, in terms of legal concepts, notable examples of inspiration from classical theoretical work include the use of principle-oriented interpretivism (Dworkin, 1988) and of a positivist account of law as analogical reasoning (Hart, 2012; Sunstein, 2018), both mapped onto Constitutional AI and case-based reasoning (Caputo, 2024). This work also employs concepts such as precedent, stare decisis (i.e., doctrine instructing the system to stand by things already decided), and case-based reasoning from common law doctrine (Llewellyn, 2016). These concepts are also formative for proposals that pair a constitution with AI courts developing AI case law (Abiri, 2024). Statutory and constitutional interpretation methods (Scalia and Garner, 2012) are also important and have been adapted as tools for resolving ambiguity in specifications (Kolt et al., 2026).

Taxonomies of Alignment Approaches. Existing surveys of alignment and governance work generally organize the field by technical methods such as RLHF (Lu et al., 2025; Shen et al., 2023), or by value-setting principles (Chalkidis, 2025). For example, one line of work distinguishes between Overton, steerable, and distributional pluralism (Sorensen et al., 2024). Others have organized the field into pluralistic, contextual, and collective mechanisms (Bao et al., 2026). Still, the field remains largely unmapped, with the various approaches emerging largely independently of each other and

not in direct conversation with one another.

Typology Construction. Motivated by taxonomies sorted by non-equivalent criteria and by free-floating typologies whose dimensions are never made explicit, prior work provides a template for building typologies from an overarching concept into explicit, clearly defined dimensions of an orthogonal nature (Collier et al., 2012). The underlying notion of a property space is due to earlier work in the same tradition (Barton, 1955), as is the distinction between conceptual and explanatory typologies (Elman, 2005). As we find equivalent failure types in existing taxonomies of alignment, we follow this methodology.

Comparison to Related Work. Existing work maps traditional concepts from philosophical, political, and legal theory onto alignment and governance methods one at a time. We take a different approach by organizing the field according to a more complete set of underlying principles by which these methods organize themselves, and map individual traditions and concepts onto them subsequently. Furthermore, we attempt to critically interrogate whether the stated inspiration in traditional work results in a faithful technical implementation, and describe ways in which gaps or confusion emerge. Additionally, this approach allows us to identify gaps in approaches to alignment and governance which remain largely unexplored in the technical literature.