

Values, ethics & morals in AI alignment

Roma Patel

[Anca Dragan]

Intro & this lesson

Research at Google DeepMind (2022-now)

[behavior, alignment]

[language grounding, multi-agent reinforcement learning]

Research in general

[safety, post-training, language understanding]


PhD at Brown University in CS (2018-2022)

[language grounding, language understanding)]


Today: What does it mean when we say we want to align language models to values? How do we currently do this and what's left?

Language model training & alignment: **the pipeline**

Language model training & alignment: the pipeline




Language model training & alignment: the pipeline



Data collection &
curation; 15T tokens,
filtering, multimodal

Language model training & alignment: the pipeline


also mid-training, continual training, annealing..



General language,
knowledge & general
capability training

Language model training & alignment


SFT, RLHF, instruction
tuning, red-teaming



Post-training
and aligning to
humans


Language model training & alignment

System prompting,
attribution,
in-context learning..




The model you use!

Language model training & alignment



When we talk about values & morals & ethics, where does it fit into this picture?

Language model training & alignment



When we talk about values & morals & ethics, where does it fit into this picture?
Where **should** it fit into the picture?

All the value-laden alignment words



Values

What matters to humans?

Formally, “*Values are things regarded as good, worthwhile, important, or worthy of pursuit.*”

happiness • autonomy • equality • truth
• loyalty • safety • beauty •
achievement • freedom

A mini timeline

Aristotle (~350 BCE)

- *eudaimonia*: human flourishing / living well
- multiple human goods and virtues

Jeremy Bentham (1789)

- pleasure / happiness as the central good

John Stuart Mill (1861)

- develops utilitarianism; famously distinguishes qualitatively different pleasures

Friedrich Nietzsche (19th c.)

- radical attention to the creation/genealogy of values

Max Scheler (early 20th c.)

- systematic hierarchy of values

Alain Locke (20th c.)

- influential pluralistic value theory; values connected to human experience

Values

Philosophy

Axiology studies what has value and why. Philosophers distinguish **intrinsic** value—valuable in itself—from **instrumental** value—valuable as a means to something else.

Different theories disagree over **whether values can be reduced to one underlying good** or **whether multiple, potentially conflicting, goods exist**.

A mini timeline

Aristotle (~350 BCE)

- *eudaimonia*: human flourishing / living well
- multiple human goods and virtues

Jeremy Bentham (1789)

- pleasure / happiness as the central good

John Stuart Mill (1861)

- develops utilitarianism; famously distinguishes qualitatively different pleasures

Friedrich Nietzsche (19th c.)

- radical attention to the creation/genealogy of values

Max Scheler (early 20th c.)

- systematic hierarchy of values

Alain Locke (20th c.)

- influential pluralistic value theory; values connected to human experience

Values

AI Alignment

Alignment systems attempt to infer or encode these values from **behavior, language, feedback, demonstrations, preferences, or explicit principles.**

safety



autonomy



Human values can refer to very different targets: **preferences, goals, principles, interests, moral beliefs, or conceptions of well-being.**

A mini timeline

Aristotle (~350 BCE)

- *eudaimonia*: human flourishing / living well
- multiple human goods and virtues

Jeremy Bentham (1789)

- pleasure / happiness as the central good

John Stuart Mill (1861)

- develops utilitarianism; famously distinguishes qualitatively different pleasures

Friedrich Nietzsche (19th c.)

- radical attention to the creation/genealogy of values

Max Scheler (early 20th c.)

- systematic hierarchy of values

Alain Locke (20th c.)

- influential pluralistic value theory; values connected to human experience

Pluralism

What if many things matter?

Value pluralism is the **view that multiple genuine values exist and may be irreducible, incommensurable, or in conflict.**

Liberty ↔ Equality

Truth ↔ Kindness

Privacy ↔ Safety

Autonomy ↔ Welfare

A mini timeline

Aristotle
multiple human goods

↓

William James / Max Weber
conflicting ideals & value spheres

↓

Isaiah Berlin
value pluralism &
incommensurability

↓

Contemporary pluralism
reasonable disagreement

↓

Pluralistic alignment
representing multiple human values

Pluralism

AI Alignment

What should an AI do when legitimate values conflict?

Alignment often **compresses judgments** into a single reward, ranking, or policy, even when the underlying values may not admit a single ordering.

Pluralistic alignment instead asks how models can **represent, preserve, and appropriately navigate multiple values and viewpoints.**

safety

autonomy



A mini timeline

Aristotle
multiple human goods



William James / Max Weber
conflicting ideals & value spheres



Isaiah Berlin
value pluralism &
incommensurability



Contemporary pluralism
reasonable disagreement



Pluralistic alignment
representing multiple human values

Morals

What is right and wrong?

“Principles, judgments, or norms concerning right and wrong conduct and how people ought to treat one another.”

Philosophers disagree! **Descriptive** morality (what a person or community actually regards as right and wrong) vs **normative** morality (what people ought to regard as right and wrong)

A mini timeline

Consequentialism
Judge actions by outcomes.

Key figures:
Bentham → Mill → Sidgwick

Deontology
Some actions are required or forbidden independently of maximizing outcomes.

Central figure:
Immanuel Kant

Virtue ethics
Instead of:
“What should I do?”
“What kind of person should I be?”

Central figure:
Aristotle

Morals

AI alignment

Moral alignment aims to safeguard against violating a user's moral assumptions.

Moral relativism holds that **moral truth**, **justification**, or **correctness** is relative to some framework—such as a **culture**, **society**, or **individual**—rather than universally invariant.

India



Japan



US



A mini timeline

Consequentialism
Judge actions by outcomes.

Key figures:
Bentham → Mill → Sidgwick

Deontology
Some actions are required or forbidden independently of maximizing outcomes.

Central figure:
Immanuel Kant

Virtue ethics
Instead of:
"What should I do?"
"What kind of person should I be?"

Central figure:
Aristotle

Ethics

Morality ≈ the object

Ethics ≈ systematic reasoning about it for purposes of this lecture.

Ethics is the systematic **study of questions about what is good, right, permissible, obligatory, and how we ought to live or act.**

[Metaethics: what even makes something morally true/right?]

Mini timeline

Ancient ethics

Socrates → Plato → Aristotle

How should we live?

↓

Modern moral theory

Hume • Kant • Bentham • Mill

What makes actions right?

↓

20th-century ethics

metaethics + normative + applied ethics

↓

Computer / machine ethics

↓

AI ethics & alignment

Normative ethics

What principles determine what we ought to do?

Consequentialism

Deontology

Virtue ethics

Contractualism

Care ethics

Applied ethics

What should we do in this particular domain?

Medical ethics

Business ethics

AI ethics

Bioethics

Environmental ethics

Ethics

AI alignment

In AI alignment, “Ethical AI” can mean at least two different things:

Ethical **behavior** (does the model **act consistently** with some **ethical principles**?)

Ethical **design** (are the **processes** used to **build/deploy** the system themselves **ethically justified**?)

Alignment frequently focuses on the former while ethics can interrogate the choice of principles itself.

Mini timeline

Ancient ethics

Socrates → Plato → Aristotle

How should we live?

↓

Modern moral theory

Hume • Kant • Bentham • Mill

What makes actions right?

↓

20th-century ethics

metaethics + normative + applied ethics

↓

Computer / machine ethics

↓

AI ethics & alignment

Normative ethics

What principles determine what we ought to do?

Consequentialism

Deontology

Virtue ethics

Contractualism

Care ethics

Applied ethics

What should we do in this particular domain?

Medical ethics

Business ethics

AI ethics

Bioethics

Environmental ethics

Norms

What does a group expect people to do?

Norms are **socially shared expectations, rules, or standards** of behavior.

Norm $\neq$ moral truth.

A society can have a widely followed norm that is morally objectionable.

A mini timeline

Hume
conventions & social practices
↓
Durkheim
social rules & collective life
↓
Bicchieri
social norms & expectations
↓
NLP datasets
Social Chemistry, Moral Stories, MIC...
↓
LLMs
learning norms from language +
feedback

Norms

AI alignment

Pretraining exposes models to enormous amounts of **implicit** social norms.

Post-training makes some norms more **explicit**:

- Be polite.
- Don't insult the user.
- Respect privacy.
- Don't provide dangerous assistance.

The alignment question becomes: Should AI **learn existing norms**, **follow** them, or sometimes **challenge** them?

A mini timeline

Hume
conventions & social practices
↓
Durkheim
social rules & collective life
↓
Bicchieri
social norms & expectations
↓
NLP datasets
Social Chemistry, Moral Stories, MIC...
↓
LLMs
learning norms from language +
feedback

Preferences

What do people want or choose?

A preference is a **comparative judgment** that one **option**, **outcome**, or **state** is favored over another.

Preference $\neq$ Value

Preference $\neq$ Moral judgment

A mini timeline

Bentham
utility

↓

Pareto
preferences & welfare

↓

von Neumann & Morgenstern
expected utility

↓

Arrow / social choice
aggregating preferences

↓

RLHF / preference learning
human comparisons $\rightarrow$ model
objectives

Preferences

Philosophy

Preferences became central to **decision theory**, **economics**, and social choice because they allow **subjective desirability** to be represented **mathematically**.

Classical rational-choice models often impose properties such as:

completeness • transitivity • consistency

to permit utility representation.

A mini timeline

Bentham
utility

↓

Pareto
preferences & welfare

↓

von Neumann & Morgenstern
expected utility

↓

Arrow / social choice
aggregating preferences

↓

RLHF / preference learning
human comparisons → model
objectives

Preferences

AI alignment

Preferences are one of the most common observable proxies for values.

Humans prefer Response A > Response B



Preference / reward model



RLHF / DPO



Model behavior

**Preference ≠ Value ≠ Moral
Judgment**

I may **prefer** cake,
value health,
and believe I **ought** to eat well.

A mini timeline

Bentham
utility



Pareto
preferences & welfare



von Neumann & Morgenstern
expected utility



Arrow / social choice
aggregating preferences



RLHF / preference learning
human comparisons → model
objectives

Synthesis

Values – What matters?

Morality – What is right/wrong?

Ethics – How do we reason about it?

Norms – What does society expect?

Pluralism – What if many things matter?

Preferences – What do I choose?

I choose it
↓
Preference
It matters to me
↓
Value
People around me expect it
↓
Norm
I think people ought to do it
↓
Moral judgment
I have a theory for why they ought to do it
↓
Ethics
Other genuine values conflict with it
↓
Pluralism
Other people disagree with me
↓
Legitimacy
We formalize the principles + decision process
↓
Constitution
↓
 Alignment

The rise of benchmarking of values, morals & ethics

First question to ask: if we take models that are deployment-ready—where do they stand on ethics? Do they reason like humans? Do they tell us that they “have” values?

we turned this into an ML benchmark question

Evaluating morality




Figure 1: Given different scenarios, models predict widespread moral sentiments. Predictions and confidences are from a BERT-base model. The top three predictions are incorrect while the bottom three are correct. The final scenario refers to Bostrom (2014)'s paperclip maximizer.

ETHICS; Hendrycks et al., ICLR 2021


Do model reason about commonsense moral problems as most humans would?

Subsets: Commonsense morality · Deontology · Justice · Utilitarianism · Virtue ethics

MoralChoice; Scherrer et al., NeurIPS 2023

Harder task: when scenarios are particularly ambiguous, context-dependent do model stances on morality change?

Truly held vs. context/user-dependent beliefs



Evaluating ethical dilemmas




Figure 1: Example narrative found in *Moral Stories*. Jenna — the **actor** — performs **actions** to fulfill her **intention** against the background of the **situation**, by either following or violating the specified social **norm**. The **consequences** describe the actions' effects on the actor and their environment.

MoralStories; Emelin et al.,
EMNLP 2021

When ethical dilemmas
are grounded in real
narrative stories does
this change model
reasoning?

MultiTP; Jin et al., ICLR 2025

Trolley problems and
multilingual trolley problems

Does language ID and cultural
information affect moral
reasoning about ethical
dilemmas & trolley problems?




Figure 1: An example scenario in the MULTITP dataset. Each question is presented in 107 different languages. Here, we select three languages, German, Chinese, and Swahili, and show the responses of LLMs (English translations provided for readers).

Evaluating norms



SocialChem101; Forbes et al., EMNLP 2020

Crowdsourced social norms from humans from different cultures, under different types of pressure/contexts

[large enough for train]

Moral Integrity; Ziems et al., CL 2022





Figure 1: A representative MIC annotation. We evaluate the AI response (**Reply**) to a human query (**Prompt**) using Rules of Thumb (**RoT**), which describe “right and wrong” ways to handle the conversation. Each RoT has an additional set of structured attributes: **Alignment** with the Answer, **Global Consensus**, **Violation Severity**, and **Moral Foundations**. There is also a **Revised Answer** that aligns with the RoT. See Table 5 for more examples.

Crowdsourced “rules of thumb” that explain moral convictions, integrity, severity of violations.

[large enough for training]

Evaluating values



Collect real human value data and measure a model's ability to simulate that user's values (based on demographics, language, Value categories)

OpinionQA; Santurkar et al., 2023

Poll humans in PEW-like settings, and then evaluate LLM opinion distributions. Can evaluate:

- political leaning
- stance on partisan topics
- opinion similarity to population groups

On a scale of 1 to 4, 1 meaning 'Very important' and 4 meaning 'Not at all important', how important is leisure time in your life?




Figure 1: Human values often depend on cultural contexts, such as, country, residential area, and education. Given a value question and demographic attributes, we examine if a language model exhibits awareness of the human answer distribution.

Ethical *behavior* instead of just correctness




Figure 2. A mock-up of a game in the MACHIAVELLI benchmark, a suite of text-based environments. At each step, the agent observes the scene and a list of possible actions; it selects an action from the list. Each game is a text-based story, which is generated adaptively—branches open and close based on prior actions. The agent receives a reward when it achieves one of the goals. By densely annotating our environment, we are able to construct a behavioral report of the agent and measure the trade-off between rewards and ethical behavior.

MACHIAVELLI; Pan et al., 2023

In open-ended tasks in carefully crafted environments can we evaluate behaviors like tendencies to be power-seeking, cause disutilities, commit ethical violations?

Jiminy Cricket; Hendrycks et al., 2022

In open-ended text-based adventure games, can we design inputs/scenarios such that LLMs would have to face ethical dilemmas we can test?





Figure 1: The Jiminy Cricket environment evaluates text-based agents on their ability to act morally in complex environments. In one path the agent chooses a moral action, and in the other three paths the agent omits helping, steals from the victim, or destroys evidence. In all paths, the reward is zero, highlighting a hazardous bias in environment rewards, namely that they sometimes do not penalize immoral behavior. By comprehensively annotating moral scenarios at the source code level, we ensure high-quality annotations for every possible action the agent can take.


The kinds of insights we can get from this

The kinds of insights we can get from this



What's the safest model?
Which should I be using?
What should I invest in?

The kinds of insights we can get from this




How good is my model at every stage? Do training methods upstream influence more safety downstream?

What's the safest model?
Which should I be using?
What should I invest in?

Integrating this into model training


Integrating this into model training



Constitutional AI & Collective CAI; Anthropic, 2022/23

CAI: Write down values, ethical principles, guidelines that models should follow → train models on this.

CCAI: Get collective, public input on guidelines



Value Steering; Sorensen et al., 2024

Explicitly steer a model to adapt to different user's value profiles instead of only representing aggregate viewpoints.




Figure 1: The value profile autoencoder setup. Decoder outputs are from trained profile decoder while demographics are illustrative to preserve privacy. The encoder extracts/compresses value information from rater examples, and the decoder changes predictions on held-out questions according to the value profile.

Top-down vs. bottom-up approaches

Case study: **top-down alignment** of agents for morality

What? Formulate the set of moral principles and how you would evaluate them

How? Then add this signal into reward functions, data filtering/weighting, inference-time processing and so on

Case study: **top-down alignment** of agents for morality

What? Formulate the set of moral principles and how you would evaluate them

How? Then add this signal into reward functions, data filtering/weighting, inference-time processing and so on

The problem? What's the set of moral principles? Who decides this? Is it exhaustive?

this is an open question
that no moral philosophers
have an answer to :)

Case study: **top-down alignment** of agents for morality

What? Formulate the set of moral principles and how you would evaluate them

How? Then add this signal into reward functions, data filtering/weighting, inference-time processing and so on

The problem? What's the set of moral principles? Who decides this? Is it exhaustive?

What happens when principles conflict?

["Don't kill humans", "Don't lie", "Don't hurt anything", "Be fair"]

Case study: **top-down alignment** of agents for morality

What? Formulate the set of moral principles and how you would evaluate them

How? Then add this signal into reward functions, data filtering/weighting, inference-time processing and so on

The problem? What's the set of moral principles? Who decides this? Is it exhaustive?

Limitations :

Hard to be prescriptive about rules that cover all situations

Moral rules conflict with each other

Different people/perspectives don't want the same thing

Case study: **bottom-up alignment** of agents for morality

Case study: **bottom-up alignment** of agents for morality

What? Don't prescriptively formulate the principles; ask humans, collect large amounts of this data, and learn from this data

How? inverse RL, imitation learning, RLHF, post-training methods

Case study: **bottom-up alignment** of agents for morality

What? Don't prescriptively formulate the principles; ask humans, and learn from this data

How? inverse RL, imitation learning, RLHF, post-training methods

The problem? Whose data do you use? What if you're not fair/representative enough?

What if people disagree (and rightly so) but our current objectives don't allow us to handle that?

Case study: **bottom-up alignment** of agents for morality

What? Don't prescriptively formulate the principles; ask humans, and learn from this data

How? inverse RL, imitation learning, RLHF, post-training methods

The problem? Whose data do you use? What if you're not fair/representative enough?

What if a person's stance changes during the conversation?

Can a model's value system update on the fly?

What approach is better?

- **No consensus** on the “correct” approach
- We don't know how to guarantee perfectly moral behavior, perfect value alignment, preserve disagreement
- But we know how to measure **better** or **worse** and hill-climb on existing metrics
- [but metrics get saturated and keep needing to be rebuilt]

What approach is actually used?

1. The most **technically feasible** one that leads to the best benchmark **performance**

RLHF,
constitutions,
aggregating
preferences,
evaluating ethics

What approach is actually used?

1. The most **technically feasible** one that leads to the best benchmark **performance**
2. Start with the easy stuff: we know what things are strict red lines (manipulation, harm to humans, mental health). So strictly encode these as **rules** the model is trained to **not violate**

RLHF,
constitutions,
aggregating
preferences,
evaluating ethics

Child safety,
graphic images,
manipulation of
humans

What approach is actually used?

1. The most **technically feasible** one that leads to the best benchmark **performance**
2. Start with the easy stuff: we know what things are strict red lines (manipulation, harm to humans, mental health). So strictly encode these as **rules** the model is trained to **not violate**
3. The **softer** things (e.g., adapting to value systems) better objectives, evaluation, training, datasets can aim to capture
4. Lots of research on **training objectives** (pre-training/fine-tuning), **dataset construction**, **attribution**, **reward design**, **behavior evaluation** to get this right


RLHF,
constitutions,
aggregating
preferences,
evaluating ethics

Child safety,
graphic images,
manipulation of
humans

Value steering,
preserving
disagreement,
deliberation,
diversity of data

So much work here
to fix the softer
questions

Where does all the current work lie



Where does all the current work lie



Humans
(create content)



The internet
(web pages, books,
forums, etc.)



Other sources
(e.g., code, datasets)

Research on data
diversity,
curating data,
cultural
differences,
global opinions,
multimodal data

Deduplicating Training Data Makes Language Models Better

Katherine Lee^{*†}

Daphne Ippolito^{*†‡}

Andrew Nystrom[†]

Chiyuan Zhang[†]

Douglas Eck[†]

Chris Callison-Burch[‡]

Nicholas Carlini[†]

Whose View of Safety? A Deep DIVE Dataset for Pluralistic Alignment of Text-to-Image Models

Charvi Rastogi
Google DeepMind

Tian Huey Teh
Google DeepMind

Pushkar Mishra
Google DeepMind

Roma Patel
Google DeepMind

Ding Wang
Google Research

Mark Diaz
Google Research

Alicia Parrish
Google DeepMind

Aida Mostafazadeh Davani
Google Research

Zoe Ashwood
Google DeepMind

Michela Paganini
Google DeepMind

Vinodkumar Prabhakaran
Google Research

Verena Rieser
Google DeepMind

Lora Aroyo
Google DeepMind

The PRISM Alignment Dataset: What Participatory, Representative and Individualised Human Feedback Reveals About the Subjective and Multicultural Alignment of Large Language Models

Hannah Rose Kirk^{1*} **Alexander Whitefield**² **Paul Röttger**³ **Andrew Bean**¹
Katerina Margatina^{4‡} **Juan Ciro**^{5,11} **Rafael Mosquera**^{5,6} **Max Bartolo**^{7,8}


Adina Williams⁹ **He He**¹⁰ **Bertie Vidgen**^{1,11†} **Scott A. Hale**^{1,12†}

¹University of Oxford ²University of Pennsylvania ³Bocconi University

⁴AWS AI Labs ⁵ML Commons ⁶Factored AI ⁷UCL ⁸Cohere

⁹MetaAI ¹⁰New York University ¹¹Contextual AI ¹²Meedan

What are the open questions in each part of the pipeline?



Research on
documenting
practices,
transparency,
ensuring
representation,
fairness

Documenting Large Webtext Corpora: A Case Study on the Colossal Clean Crawled Corpus


Jesse Dodge♣ Maarten Sap♣♥ Ana Marasović♣♥ William Agnew♥
Gabriel Ilharco♥ Dirk Groeneveld♣ Margaret Mitchell♣ Matt Gardner♣
♥Paul G. Allen School of Computer Science & Engineering, University of Washington
♣Hugging Face
♣Allen Institute for Artificial Intelligence
jessed@allenai.org

Data Statements for Natural Language Processing: Toward Mitigating System Bias and Enabling Better Science

Emily M. Bender
Department of Linguistics
University of Washington
ebender@uw.edu

Batya Friedman
The Information School
University of Washington
batya@uw.edu

What are the open questions in each part of the pipeline?



What are the open questions in each part of the pipeline?

Fine-Tuning Language Models from Human Preferences

Daniel M. Ziegler* Nisan Stiennon* Jeffrey Wu Tom B. Brown

Alec Radford Dario Amodei Paul Christiano Geoffrey Irving

OpenAI

{dmz,nisan,jeffwu,tom,alec,damodei,paul,irving}@openai.com

Collective Constitutional AI: Aligning a Language Model with Public Input

Saffron Huang[†]
saffron@deepmind.org
Collective Intelligence Project
San Francisco, California, USA

Divya Siddharth[†]
divys@deepmind.org
Collective Intelligence Project
San Francisco, California, USA

Liane Lovitt[†]
Anthropic
San Francisco, California, USA

Thomas I. Liao[‡]
Anthropic
San Francisco, California, USA

Esin Durmus
Anthropic
San Francisco, California, USA

Alex Tamkin
Anthropic
San Francisco, California, USA

Deep Ganguli[†]
Anthropic
San Francisco, California, USA
deep@anthropic.com

Value Profiles for Encoding Human Variation

Isen^{1†}, Pushkar Mishra², Roma Patel², Michael Henry T
Georgina Evans², Iason Gabriel², Noah Goodman², Ve
of Computer Science, University of Washington, Seattle, W

²Google DeepMind, London, UK

[†]Work done during an internship at Google DeepMind.


[†]13@cs.washington.edu, verenarieser@deepmind.com

PERSONALIZED SOUPS: PERSONALIZED LARGE LANGUAGE MODEL ALIGNMENT VIA POST-HOC PARAMETER MERGING

Joel Jang^{1,2} Seungone Kim³ Bill Yuchen Lin² Yizhong Wang¹ Jack Hessel²
Luke Zettlemoyer¹ Hannaneh Hajishirzi^{1,2} Yejin Choi^{1,2} Prithviraj Ammanabrolu⁴

¹University of Washington ²Allen Institute for AI ³KAIST AI ⁴UC San Diego

joeljang@cs.washington.edu



Research on actual post-training alignment and methods that allow for safest models

What are the open questions in each part of the pipeline?

Whose Opinions Do Language Models Reflect?

Shibani Santurkar
Stanford
shibani@stanford.edu

Esin Durmus
Stanford
esindurmus@cs.stanford.edu

Faisal Ladhak
Columbia University
faisal@cs.columbia.edu

Cinoo Lee
Stanford
cinoollee@stanford.edu

Percy Liang
Stanford
pliang@cs.stanford.edu

Tatsunori Hashimoto
Stanford
thashim@stanford.edu

Societal Impacts

Towards measuring the representation of subjective global opinions in language models



CultureBank: An Online Community-Driven Knowledge Base Towards Culturally Aware Language Technologies

Weiyan Shi[†], Ryan Li[△], Yutong Zhang[△], Caleb Ziems[△], Chunhua Yu[△],
Raya Horesh[◇], Rogério Abreu de Paula[◇], Diyi Yang[△]
Northeastern University[†], Stanford University[△], IBM Research[◇]
{lansong,yutongz7,cziems,syu03}@stanford.edu
us.ibm.com, ropaula@br.ibm.com, diyiy@stanford.edu

SOCIAL CHEMISTRY 101:
Learning to Reason about Social and Moral Norms

Maxwell Forbes^{†‡} Jena D. Hwang[†] Vered Shwartz^{†‡} Maarten Sap[†] Yejin Choi^{†‡}

[†]Paul G. Allen School of Computer Science & Engineering, University of Washington

[‡]Allen Institute for AI

{mbforbes,msap,yejin}@cs.washington.edu, {jenah,vereds}@allenai.org

maxwellforbes.com/social-chemistry

4. Deployed model


Interact with users
in real-world contexts




End user
(asks questions,
gets answers,
uses the model
in their life)

Research on
inference-time
updates,
interpretability,
evaluation


What are the open questions in each part of the pipeline?



What are the open questions in each part of the pipeline?




What are the open questions in each part of the pipeline?



Who do we collect from? How do we ensure representation, diversity, conflicting data?

How do we input this data to models? Do we downweight or upweight for representation? Better formats of value reps?

What are the open questions in each part of the pipeline?




Who do we collect from? How do we ensure representation, diversity, conflicting data?


How do we input this data to models? Do we downweight or upweight for representation? Better formats of value reps?

What pre-training objective can handle this to best help post-training?


What are the open questions in each part of the pipeline?



What are the open questions in each part of the pipeline?



What are the open questions in each part of the pipeline?



Philosophy <--> Computation



<< Look, man, we haven't got all day. How about Jim and I just figure it out between ourselves? >>

<https://www.evilaicartoons.com/>

A lot of these tricky, philosophical questions are useful now

Can be materialized into:

- Measurement & metrics
- Math & objectives
- Dataset curation
- Policy guidelines

Is necessary for:

- Foresight
- Risk assessment
- Interpreting models

Thank you! Questions?

Extra Slides

How much data is needed?

Metric	Estimate (Human)	Equivalent in Tokens
Reading Speed	~250 words per minute	~333 tokens per minute
Daily Consumption	4 hours of reading/heavy media consumption a day	~80,000 tokens per day
Yearly Consumption	365 days of active consumption	~29 million tokens per year
Lifetime Total (80 years)	A lifetime of constant learning and reading	~2.33 billion tokens

Pretraining Data Stats by Model Generation

Model Version	Estimated Dataset Size (Raw Data)	Estimated Token Count
GPT-3	~45 Terabytes (unfiltered) / 570 GB (filtered)	300 Billion tokens
GPT-4	~1 Petabyte	~13 Trillion tokens (including multiple epochs)
GPT-4o / GPT-5 Series	Multi-petabyte (incorporates massive audio/visual data)	Estimated 15+ Trillion tokens



<< Look, man, we haven't got all day. How about Jim and I just figure it out between ourselves? >>

<https://www.evilaicartoons.com/>

Relativism

Does what's right depend on who/where you are?

Moral relativism holds that **moral truth**, **justification**, or **correctness** is relative to some framework—such as a **culture**, **society**, or **individual**—rather than universally invariant.

A mini timeline

Protagoras
"Man is the measure"

↓

Cultural anthropology
Boas → Benedict → Herskovits
cultural variation in norms

↓

Gilbert Harman
moral relativism

↓

David Wong
pluralistic relativism

↓

AI alignment
Should models adapt morally across
cultures?

Relativism

AI alignment

Models operate across users and cultures with different norms and values.

Should a model:

- follow universal principles?
- adapt to local cultural norms? adapt to user?
- combine universal constraints with local variation?

India



Japan



US



A mini timeline

Protagoras
"Man is the measure"

↓

Cultural anthropology
Boas → Benedict → Herskovits
cultural variation in norms

↓

Gilbert Harman
moral relativism

↓

David Wong
pluralistic relativism

↓

AI alignment
Should models adapt morally across cultures?